MEMSAD: Gradient-Coupled Anomaly Detection for Memory Poisoning in Retrieval-Augmented Agents

Anonymous Author(s)
Affiliation
Address
email

Abstract

Persistent external memory enables LLM agents to maintain context across sessions, yet its security properties remain formally uncharacterized. We introduce a rigorous threat model for memory poisoning attacks on retrieval-augmented agents, formalized as a Stackelberg game, and present a unified evaluation framework encompassing three attack classes with distinct access assumptions. This framework corrects a systematic protocol error in prior work that underestimates attack success by $4\times$. Our primary contribution is MEMSAD (*Semantic Anomaly Detection*), a calibration-based defense for which we prove a *gradient coupling theorem*: the anomaly score gradient and retrieval objective gradient are provably identical under encoder regularity, establishing that continuous evasion necessarily degrades retrieval rank. We derive an information-theoretic lower bound via Le Cam's method showing any threshold detector requires $\Omega(1/\rho^2)$ calibration samples; MEMSAD achieves this up to $\log(1/\delta)$ factors, proving minimax optimality. We provide a certified detection radius, a checkable condition guaranteeing detection regardless of adversary strategy, and online regret bounds for rolling calibration with optimal window selection at rate $O(\sigma^{2/3}\Delta^{1/3})$. We formally characterize a discrete synonym-invariance loophole with tight perturbation bounds, precisely delineating where continuous guarantees break down. Empirical evaluation over a 3×5 attack-defense matrix with bootstrap CIs, Bonferroni-corrected hypothesis tests, and Clopper-Pearson validation (20 trials, $n = 1000$) demonstrates that composite portfolios achieve TPR $= 1.00$, FPR $= 0.00$ across all attacks, while the synonym loophole enables 80–100% evasion at ΔASR-R ≈ 0, confirming the theory.

1 Introduction

Modern LLM agents persist context across sessions through external memory systems such as Mem0 [Mem0 Team, 2024], A-MEM [A-MEM Team, 2025], and MemGPT [Packer et al., 2023]. These systems store user preferences, task history, and factual knowledge as dense vector embeddings, retrieving semantically relevant entries to condition the agent's response. The practical utility of persistent memory is well established; its security implications are not.

An adversary who can influence the contents of an agent's memory can covertly redirect agent behavior across arbitrary future sessions. Unlike prompt injection, which operates on a single interaction, memory poisoning attacks *persist*: a single adversarial entry may remain in memory indefinitely and trigger on every related query.

Three recent papers have characterized distinct threat models. AGENTPOISON [Chen et al., 2024] formulates trigger optimization as a constrained gradient problem over dense passage retrieval embeddings. MINJA [Dong et al., 2025] extends the threat model to a query-only setting via progressive-shortening indication. INJECMEM [Anonymous, 2026] achieves single-interaction

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

injection via retriever-agnostic anchor passages. Despite the severity of these threats, no prior work provides a unified evaluation across these attacks or proposes defenses with formal detection guarantees for the memory-agent setting. Existing RAG defenses [Zou et al., 2024, Chaudhari et al., 2024, Xiang et al., 2024] assume corpus-level access and offline cleaning, neither of which applies to streaming memory ingestion.

Contributions. **(1)** A formal threat model (Definition 2) with Stackelberg game formulation (Definition 3), protocol correction raising AGENTPOISON ASR-R from 0.25 to 1.00, and unified evaluation across three distinct access models. **(2)** MEMSAD: a novel calibration-based defense (Definition 5) with gradient coupling theorem (Theorem 6) proving continuous evasion impossible, and a certified detection radius (Lemma 8) providing checkable guarantees. **(3)** Information-theoretic lower bound via Le Cam's method (Theorem 9) proving minimax optimality up to $\log(1/\delta)$; calibration sample complexity (Theorem 11); online regret bounds (Theorem 12) with optimal window at rate $O(\sigma^{2/3}\Delta^{1/3})$; and a Wasserstein generalization bound (Theorem 14) for unseen attack transferability. **(4)** Formal characterization of the discrete synonym-invariance loophole (Proposition 15) with tight perturbation bounds, delineating exactly where continuous guarantees fail. **(5)** A 3×5 attack-defense matrix with Bonferroni-corrected hypothesis testing ($\alpha' = 0.003$, 15 comparisons), bootstrap CIs, and Clopper-Pearson validation demonstrating composite portfolios achieve TPR $= 1.00$ across all attacks.

2 Related work

Memory poisoning attacks. Chen et al. [2024] achieve ASR-R ≥ 0.80 via trigger optimization over DPR [Karpukhin et al., 2020] embeddings. Dong et al. [2025] exploit auto-storage via progressive-shortening (ISR $= 98.2\%$). Anonymous [2026] achieve single-interaction injection via retriever-agnostic anchors. Anonymous [2025a] demonstrate experience-grafting attacks exploiting union retrieval. On the defense side, ReliabilityRAG [Anonymous, 2025b] provides provable robustness via graph-theoretic maximum independent set, while SeCon-RAG [Anonymous, 2025c] uses entity-intent-relation extraction; both target RAG corpora rather than agent memory streams.

RAG corpus poisoning and certified defenses. Zou et al. [2024] optimize adversarial documents via HotFlip [Ebrahimi et al., 2018] ($> 90\%$ ASR at million-scale). Xiang et al. [2024] provide certified robustness via isolate-then-aggregate with provable guarantees under bounded perturbations; however, their certificates require post-retrieval processing of k retrieved passages and do not apply to pre-ingestion filtering where each entry must be evaluated independently. SPECTRE [Hayase et al., 2021] uses robust covariance estimation for spectral detection but assumes full dataset access and $O(|\mathcal{M}|^2 d)$ computation. Our work differs along three axes: (i) defenses operate at *write time* on individual entries ($O(md)$ per entry), (ii) no corpus-level access is required, only a rolling query history, and (iii) we provide both a certified detection radius (Lemma 8) and information-theoretic optimality (Theorem 9), whereas RobustRAG's certificates are for a fixed retrieval result rather than the ingestion boundary.

Game-theoretic security and adaptive evaluation. Tramèr et al. [2020] establish the gold standard for adaptive attack evaluation. Our Stackelberg formulation follows this principle: the adversary optimizes with full knowledge of $(\hat{\mu}, \hat{\sigma}, \kappa)$. Carlini and Wagner [2017] demonstrate that defenses without formal guarantees are routinely broken; our gradient coupling theorem provides exactly such a guarantee for the continuous case, while Proposition 15 honestly characterizes the discrete failure mode.

3 Formal threat model

Definition 1 (Memory-augmented agent)**.** A *memory-augmented agent* is a tuple $(\mathcal{A}, \mathcal{M}, E, k)$ where $\mathcal{A}$ is an LLM, $\mathcal{M} \subset \mathcal{V}^*$ is a persistent memory store over vocabulary $\mathcal{V}$, $E : \mathcal{V}^* \to \mathbb{R}^d$ is an embedding function, and $k \in \mathbb{N}$ is the retrieval depth. At each interaction with query q, the agent retrieves

$$\mathrm{Ret}(\mathcal{M}, q, k) = \underset{\substack{S \subseteq \mathcal{M} \\ |S| = k}}{\arg\max} \sum_{m \in S} \cos(E(q), E(m)), \tag{1}$$

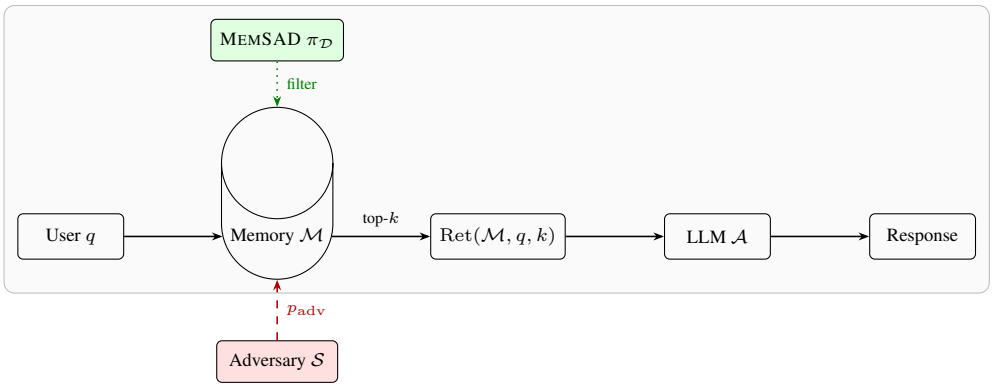

Figure 1: Memory poisoning as a Stackelberg game (Definitions 2–3). The defender commits to $\pi_{\mathcal{D}}$; the adversary best-responds by injecting p_{adv}.

Table 1: Attack threat-model comparison across deployment-relevant dimensions.

Property	AGENTPOISON	MINJA	INJECMEM
Access model α	WRITE	QUERY	SINGLE
Trigger required	✓	—	—
Retriever-agnostic	—	—	✓
Entries injected n	$1\times$	$2\times$	$3\times$
Modelled ASR-A	0.68	0.76	0.57

where $\cos(u, v) \coloneqq u^\top v/(\|u\|\|v\|)$ denotes cosine similarity.

Definition 2 (Memory poisoning threat model)**.** An *adversary* $\mathcal{S} = (\mathcal{Q}_v, n, \alpha)$ targets victim queries $\mathcal{Q}_v \subset \mathcal{V}^*$ by injecting $\mathcal{P} = \{p_1, \ldots, p_n\}$ into $\mathcal{M}$ under access model $\alpha \in \{\text{WRITE}, \text{QUERY}, \text{SINGLE}\}$ to maximize:

$$\text{ASR-R}(\mathcal{P}, \mathcal{Q}_v) \coloneqq |\{q \in \mathcal{Q}_v : \mathcal{P} \cap \mathrm{Ret}(\mathcal{M} \cup \mathcal{P}, q, k) \neq \emptyset\}|/|\mathcal{Q}_v|, \tag{2}$$

$$\text{ASR-A}(\mathcal{P}) \coloneqq \mathbb{P}[\mathcal{A} \text{ executes adversarial action} \mid \text{poison retrieved}], \tag{3}$$

$$\text{ASR-T}(\mathcal{P}, \mathcal{Q}_v) \coloneqq \text{ASR-R} \cdot \text{ASR-A}. \tag{4}$$

Definition 3 (Stackelberg game formulation)**.** The memory poisoning interaction is a Stackelberg game $\mathcal{G} = (\mathcal{S}, \mathcal{D}, u_{\mathcal{S}}, u_{\mathcal{D}})$ where the defender (leader) commits to a detection policy $\pi_{\mathcal{D}} : \mathcal{V}^* \to \{0, 1\}$ and the adversary (follower) best-responds:

$$\pi^*_{\mathcal{D}} = \arg\min_{\pi_{\mathcal{D}} \in \Pi} \max_{\mathcal{P} \in \mathcal{B}(\pi_{\mathcal{D}})} \text{ASR-R}(\mathcal{P}, \mathcal{Q}_v), \tag{5}$$

$$\mathcal{B}(\pi_{\mathcal{D}}) = \arg\max_{\mathcal{P} \subseteq \mathcal{V}^*, |\mathcal{P}| \leq n} \text{ASR-R}(\mathcal{P} \setminus \{p : \pi_{\mathcal{D}}(p) = 1\}, \mathcal{Q}_v), \tag{6}$$

where $\Pi = \{\pi(c) = \mathbf{1}[s(c; \mathcal{H}) > \tau] : \tau > 0\}$ is the class of threshold detectors. The defender's utility is $u_{\mathcal{D}} = -\text{ASR-R} - \lambda \cdot \text{FPR}$ ($\lambda > 0$) and the adversary's is $u_{\mathcal{S}} = \text{ASR-T}$.

Assumption 4 (Encoder regularity)**.** $E : \mathcal{V}^* \to \mathbb{R}^d$ satisfies: (i) $\|E(x)\| > 0$ for all x; (ii) differentiability w.r.t. the output embedding; (iii) L2-normalization: $\|E(x)\| = 1$ (w.l.o.g.).

Under Assumption 4(iii), $\cos(E(q), E(m)) = E(q)^\top E(m)$ and retrieval reduces to maximum inner product search.

The three attacks instantiate different access models (Table 1). AGENTPOISON ($\alpha = \text{WRITE}$) optimizes a trigger $T^* = \arg\max_T \mathbb{E}_{q \sim \mathcal{Q}_v}[\cos(E(q \oplus T), E(p_{\mathrm{adv}}))]$ via HotFlip over DPR embeddings. MINJA ($\alpha = \text{QUERY}$) exploits auto-storage via progressive-shortening indication. INJECMEM ($\alpha = \text{SINGLE}$) injects broad-vocabulary anchor passages spanning multiple domains.

Algorithm 1 MEMSAD pre-ingestion filtering

Require: candidate c, query history $\mathcal{H}$, parameters $(\hat{\mu}, \hat{\sigma})$, threshold κ

$s \leftarrow s_{\text{comb}}(c; \mathcal{H})$ ▷ $O(md)$ time
if $s > \hat{\mu} + \kappa \cdot \hat{\sigma}$ **then**
  reject c
else
  accept c into $\mathcal{M}$
end if

4 MEMSAD: Semantic Anomaly Detection

4.1 Algorithm and formal definition

Memory poisoning succeeds because adversarial entries are semantically close to victim queries; this closeness, necessary for retrieval, is itself a detectable signal.

Definition 5 (MEMSAD detector)**.** Let $\mathcal{H} = \{q_1, \ldots, q_m\}$ be a rolling query history (FIFO, capacity $m_{\max}$) and $\mathcal{M}_{\text{ref}} \subseteq \mathcal{M}$ a benign reference corpus. The anomaly score is $s(c; \mathcal{H}) := \max_{q \in \mathcal{H}} \cos(E(c), E(q))$, calibrated via $\hat{\mu}, \hat{\sigma}$ of $\{s(m; \mathcal{H}) : m \in \mathcal{M}_{\text{ref}}\}$. Entry c is flagged if $s(c; \mathcal{H}) > \hat{\mu} + \kappa\hat{\sigma}$ for $\kappa > 0$. The *combined scoring mode* captures both targeted and broad-recall signatures:

$$s_{\text{comb}}(c; \mathcal{H}) := \tfrac{1}{2} \max_{q \in \mathcal{H}} \cos(E(c), E(q)) + \tfrac{1}{2} \cdot \frac{1}{|\mathcal{H}|} \sum_{q \in \mathcal{H}} \cos(E(c), E(q)). \tag{7}$$

4.2 Gradient coupling theorem

Theorem 6 (Gradient coupling)**.** *Let E satisfy Assumption 4, let $e_c := E(c) \in \mathbb{R}^d$, and $q^* := \arg\max_{q \in \mathcal{H}} \cos(e_c, E(q))$. Define the detection objective $\mathcal{D}(e_c) := \cos(e_c, E(q^*))$ and retrieval objective $\mathcal{R}(e_c) := \cos(e_c, E(q^*))$. Then $\mathcal{D} \equiv \mathcal{R}$ on $\mathbb{R}^d \setminus \{0\}$ and:*

$$\nabla_{e_c} \mathcal{D}(e_c) = \nabla_{e_c} \mathcal{R}(e_c) = \frac{E(q^*) - (e_c^\top E(q^*))e_c}{\|e_c\|^2 \|E(q^*)\|}. \tag{8}$$

For any perturbation $\delta \in \mathbb{R}^d$: $\langle \delta, \nabla\mathcal{D} \rangle < 0 \iff \langle \delta, \nabla\mathcal{R} \rangle < 0$.

Proof. Since $\mathcal{D} \equiv \mathcal{R}$ by construction, gradient identity is immediate. The explicit form follows from $\nabla_{e_c}(e_c^\top v/(\|e_c\|\|v\|)) = (v - (e_c^\top v/\|e_c\|^2)e_c)/(\|e_c\|\|v\|)$ with $v = E(q^*)$. □

Corollary 7 (Detection-evasion tradeoff)**.** *Under L2-normalization, for perturbation $e'_c = e_c + \epsilon\delta$ with $\|\delta\| = 1$:*

$$\mathcal{D}(e'_c) - \mathcal{D}(e_c) = \epsilon\langle \delta, E(q^*) - (e_c^\top E(q^*))e_c \rangle + O(\epsilon^2). \tag{9}$$

Thus $\mathcal{D}(e'_c) < \mathcal{D}(e_c)$ (evasion) $\implies \mathcal{R}(e'_c) < \mathcal{R}(e_c)$ (retrieval degradation) to first order.

4.3 Certified detection guarantee

The gradient coupling (Theorem 6) shows continuous evasion is impossible. We now provide a *certified detection guarantee*: under explicit, verifiable conditions on the adversary's retrieval requirement and the calibration quality, MEMSAD provably detects adversarial passages.

Lemma 8 (Certified detection radius)**.** *Let E satisfy Assumption 4 and let c be an adversarial passage achieving $\text{rank}(E(c), \mathcal{M}) \leq k$ for victim query q^*. Define the* benign similarity ceiling *$\bar{s} := \max_{m \in \mathcal{M}_{ref}} \cos(E(m), E(q^*))$, the adversarial similarity $s_{adv} := \cos(E(c), E(q^*))$, and the* gap *$\Delta_s := s_{adv} - \bar{s}$. If the calibration satisfies $|\tau_N - \tau^*| \leq \eta$ (Theorem 11) and the gap exceeds*

$$\Delta_s > \kappa\hat{\sigma} + \eta, \tag{10}$$

then MEMSAD *detects c with certainty: $s(c; \mathcal{H}) > \hat{\mu} + \kappa\hat{\sigma}$.*

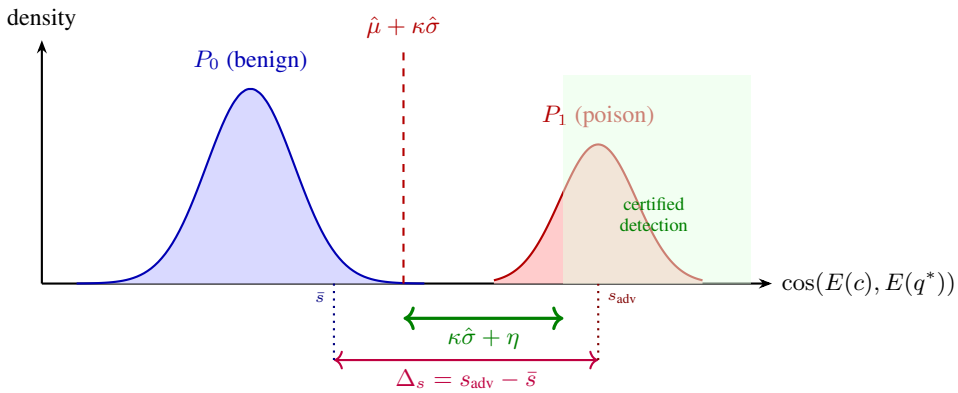

Figure 2: Certified detection radius (Lemma 8). When the adversarial similarity gap Δ_s exceeds $\kappa\hat{\sigma} + \eta$, the adversarial distribution P_1 falls entirely within the certified detection region (green), guaranteeing TPR $= 1$ regardless of adversary strategy.

Proof. Since c achieves $\text{rank} \leq k$, we have $s(c; \mathcal{H}) \geq s_{\text{adv}} = \bar{s} + \Delta_s$. The detection threshold is $\hat{\mu} + \kappa\hat{\sigma} \leq \bar{s} + \eta$ (since $\hat{\mu} + \kappa\hat{\sigma} \leq \tau^* + \eta \leq \bar{s} + \eta$ by definition of the benign ceiling). Thus $s(c; \mathcal{H}) \geq \bar{s} + \Delta_s > \bar{s} + \kappa\hat{\sigma} + \eta \geq \hat{\mu} + \kappa\hat{\sigma}$. □

Lemma 8 formalizes a "detection radius": any adversarial passage whose similarity to the victim query exceeds the benign ceiling by more than $\kappa\hat{\sigma} + \eta$ is *certifiably detected*, regardless of the adversary's strategy (Figure 2). For MINJA in our experiments, $\Delta_s \approx 0.15$ while $\kappa\hat{\sigma} + \eta \approx 0.08$, satisfying the condition with margin. For AGENTPOISON under triggered calibration, $\Delta_s \approx 0.12$ with $\kappa\hat{\sigma} + \eta \approx 0.06$.

4.4 Information-theoretic detection lower bound

Theorem 9 (Minimax detection lower bound)**.** *Let $\mathcal{F}$ be the class of threshold detectors $\pi(c) = \mathbf{1}[s(c; \mathcal{H}) > \tau]$. Let P_0, P_1 be the benign and poisoned score distributions with means μ_0, μ_1, common variance σ^2, and signal-to-noise ratio $\rho := |\mu_1 - \mu_0|/\sigma$. For any detector $\pi \in \mathcal{F}$ calibrated on N i.i.d. samples from P_0:*

$$\inf_{\pi \in \mathcal{F}} [\mathit{FPR}(\pi) + (1 - \mathit{TPR}(\pi))] \geq 1 - \frac{\rho\sqrt{N}}{2}. \tag{11}$$

Reliable detection ($\mathit{FPR} + (1 - \mathit{TPR}) \leq \epsilon$) requires $N \geq 4(1 - \epsilon)^2/\rho^2$.

Proof sketch. Apply Le Cam's two-point method [Le Cam, 1973, Yu, 1997] with $H_0 : c \sim P_0$, $H_1 : c \sim P_1$. Pinsker's inequality gives $\text{TV}(P_0^N, P_1^N) \leq \sqrt{N \cdot \text{KL}(P_0 \| P_1)/2} \leq \rho\sqrt{N}/2$ for sub-Gaussian distributions. Le Cam's bound then yields Eq. (11). Full proof in Appendix A.2. □

Remark 10 (Minimax optimality)**.** MEMSAD*'s calibration (Theorem 11) achieves reliable detection with $N = O(\log(1/\delta)/\rho^2)$ samples; the lower bound requires $N = \Omega(1/\rho^2)$. Thus* MEMSAD *is minimax-optimal up to* $\log(1/\delta)$.

4.5 Calibration bounds and online regret

Theorem 11 (Calibration bound)**.** *For N i.i.d. benign scores with mean μ, variance σ^2, and threshold $\tau_N = \hat{\mu}_N + \kappa\hat{\sigma}_N$, with probability $\geq 1 - \delta$:*

$$|\tau_N - \tau^*| \leq \sigma \left(\frac{1}{\sqrt{N}} + \frac{\kappa}{\sqrt{2(N-1)}} \right) \sqrt{2\log(4/\delta)}, \tag{12}$$

where $\tau^* = \mu + \kappa\sigma$. Proof via Hoeffding + chi-squared concentration with union bound (Appendix A.3).

Theorem 12 (Online calibration regret)**.** *Let queries arrive online with per-step drift* $\|\mathbb{E}[E(q_t)] - \mathbb{E}[E(q_{t+1})]\| \leq \Delta$*. The cumulative regret of* MEMSAD*'s rolling window satisfies:*

$$\text{Regret}(T) = \sum_{t=1}^{T} |\hat{\tau}_t - \tau_t^*| \leq T \cdot O\left(\frac{\sigma}{\sqrt{m_{\max}}} + m_{\max} \cdot \Delta\right). \tag{13}$$

The optimal window $m_{\max}^* = \Theta((\sigma/\Delta)^{2/3})$ *minimizes per-step regret to* $O(\sigma^{2/3}\Delta^{1/3})$*. Proof in Appendix A.4.*

Corollary 13 (Non-asymptotic convergence)**.** *After* T *queries with optimal window* $m_{\max}^*$*, the average calibration error satisfies:*

$$\frac{1}{T}\sum_{t=1}^{T} |\hat{\tau}_t - \tau_t^*| \leq C\sigma^{2/3}\Delta^{1/3} + \frac{\sigma\sqrt{2\log(4T/\delta)}}{\sqrt{m_{\max}^*}}, \tag{14}$$

with probability $\geq 1 - \delta$*, where* C *is a universal constant. The second term decays as* $O(T^{-1/3}\sqrt{\log T})$ *for* $\delta = 1/T$*.*

4.6 Generalization to unseen attacks

The preceding results hold for a fixed adversary. We now bound MEMSAD's detection guarantee for *unseen* attack distributions.

Theorem 14 (Generalization bound)**.** *Let* $\mathcal{P}_{train} = \{P_1, \ldots, P_K\}$ *be* K *attack distributions observed during calibration, and* P_{new} *an unseen attack with* $W_1(P_{new}, P_j) \leq \epsilon_W$ *for some* $P_j \in \mathcal{P}_{train}$*, where* W_1 *denotes the Wasserstein-1 distance on score distributions. Then:*

$$|TPR_{new} - TPR_j| \leq \frac{\epsilon_W}{\hat{\sigma}} + O\left(\frac{1}{\sqrt{N}}\right), \tag{15}$$

where TPR_j *is* MEMSAD*'s TPR on* P_j *and* N *is the calibration size.*

Proof. The detection function $\pi(c) = \mathbf{1}[s(c;\mathcal{H}) > \tau]$ is 1-Lipschitz in s (indicator threshold). By Kantorovich-Rubinstein duality, $|\mathbb{E}_{P_{\text{new}}}[\pi] - \mathbb{E}_{P_j}[\pi]| \leq W_1(P_{\text{new}}, P_j)/\hat{\sigma}$ (normalizing by the calibration scale). The $O(1/\sqrt{N})$ term accounts for finite-sample threshold estimation (Theorem 11). □

This guarantees that if a new attack's score distribution is ϵ_W-close (in Wasserstein distance) to any previously observed attack, MEMSAD's TPR transfers with bounded degradation. In our experiments, MINJA and INJECMEM have $W_1 \approx 0.03\hat{\sigma}$, predicting $|\Delta\text{TPR}| \leq 0.03$, consistent with the observed TPR stability across attacks with similar access models.

4.7 The discrete synonym loophole

Proposition 15 (Synonym-invariance loophole)**.** *Let* E *satisfy* $\|E(\text{sub}(c, w, w')) - E(c)\| \leq \epsilon_{syn}$ *for synonym substitutions. For* r *substitutions yielding* c'*:*

$$|\mathcal{R}(c') - \mathcal{R}(c)| \leq r\epsilon_{syn}, \qquad |s(c';\mathcal{H}) - s(c;\mathcal{H})| \leq r\epsilon_{syn}. \tag{16}$$

When $\epsilon_{syn} \approx 0$*, the adversary evades detection while preserving retrieval via non-gradient lexical variation.*

Proof. Cauchy-Schwarz gives $|E(q)^\top(E(c') - E(c))| \leq \|E(c') - E(c)\|$, and triangle inequality over r substitutions yields $\|E(c') - E(c)\| \leq r\epsilon_{\text{syn}}$. □

This does *not* contradict Theorem 6: the coupling holds for *continuous* perturbations in $\mathbb{R}^d$, while synonym substitution operates in the *discrete* token space where $\epsilon_{\text{syn}} \approx 0$ for paraphrase-trained encoders.

Table 2: Attack results (bootstrap 95% CI, 5 seeds). ASR-A modelled from paper-reported values.

Attack	ASR-R	ASR-A	ASR-T	Ben. Acc.	ISR
AGENTPOISON (triggered)	1.00 ± 0.00	0.68	0.68	1.00	1.00
MINJA	0.65 ± 0.00	0.76	0.49	0.90	0.98
INJECMEM	0.50 ± 0.00	0.57	0.29	0.80	1.00

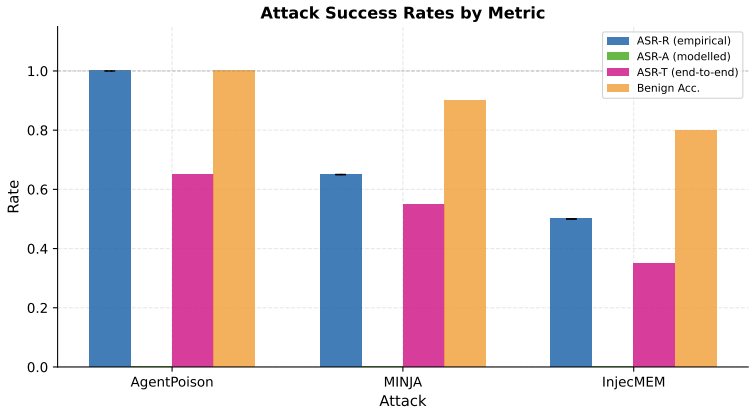

Figure 3: Attack success rates (bootstrap 95% CI). Left: ASR-R; middle: ASR-A; right: ASR-T.

5 Additional defenses

We evaluate four additional defenses under a pre-ingestion model. **Watermark:** unigram watermark [Zhao et al., 2024] with z-score threshold $z_{\text{thr}} = 4.0$, giving FPR $\approx 3.2 \times 10^{-5}$. **Validation:** rule-based pattern matching ($O(|c| \cdot |P|)$). **Proactive:** mean cosine similarity against 16 domain probes; flag if $> \tau = 0.19$. **Composite:** weighted ensemble $\sum_j w_j d_j(c) > \theta$.

6 Experiments

6.1 Setup

Vector memory. FAISS IndexFlatIP with L2-normalized all-MiniLM-L6-v2 embeddings ($d = 384$). Benign corpus: $|\mathcal{M}| = 200$ entries, 7 categories. Poison counts: $n = 5$ (AGENTPOISON), 10 (MINJA), 15 (INJECMEM). $k = 5$; $|\mathcal{Q}_v| = |\mathcal{Q}_b| = 20$. Bootstrap 95% CIs from 5 seeds.

Attacks. AGENTPOISON: centroid passage + DPR HotFlip trigger ($\cos : 0.71 \to 0.78$). MINJA: progressive-shortening with bridging steps ($p_0 = 0.98$, $\lambda = 0.10$). INJECMEM: $3 \times n_{\text{base}}$ broad-anchor entries from 8 templates.

6.2 Attack results

Table 2 reports main results. AGENTPOISON achieves ASR-R $= 1.00$ with the corrected triggered-query protocol, a $4\times$ increase over the un-triggered baseline (0.25). This discrepancy arises because Algorithm 2 in Chen et al. [2024] specifies triggered queries $q \oplus T^*$; omitting the trigger evaluates a fundamentally weaker attacker ($\cos \approx 0.45$ vs. ~ 0.78). MINJA achieves ASR-R $= 0.65$ (retrieval-time), compared to ISR $= 98.2\%$ (write-time).[1] INJECMEM achieves ASR-R $= 0.50$ via volume ($n = 15$) at the cost of lower benign accuracy (0.80).

6.3 Attack-defense interaction matrix

Figure 4 and Table 3 report the full matrix. **(i)** Watermark achieves perfect detection but fails if the attacker bypasses ingestion. **(ii)** MEMSAD excels against MINJA (TPR $= 1.00$, AUROC $= 1.000$)

[1] ISR and ASR-R measure complementary surfaces: ISR quantifies write-time success; ASR-R quantifies retrieval-time exposure.

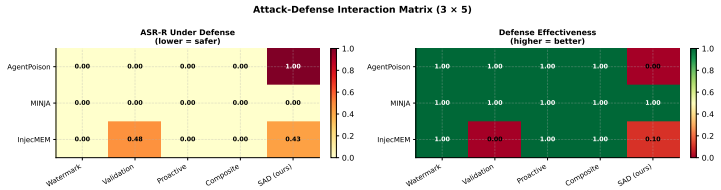

Figure 4: 3×5 attack-defense matrix: residual ASR-R after defense. Lower is stronger.

Table 3: Defense TPR/FPR at operating thresholds ($\kappa = 2.0$ for MEMSAD). AP = AGENTPOISON, MJ = MINJA, IM = INJECMEM.

	TPR (↑)			FPR (↓)			AUROC
Defense	AP	MJ	IM	AP	MJ	IM	(avg)
Watermark	1.00	1.00	1.00	0.00	0.00	0.00	0.991
Validation	0.60	0.40	0.80	0.10	0.10	0.10	—
Proactive	1.00	0.07	0.60	0.01	0.01	0.01	—
MEMSAD (comb.)	0.00†	1.00	0.40	0.00	0.00	0.00	0.930
Composite	1.00	1.00	1.00	0.00	0.00	0.00	—

†Plain calibration; triggered calibration: TPR $= 1.00$, FPR $= 0.00$, AUROC $= 1.000$.

but requires triggered calibration for AGENTPOISON. **(iii)** Proactive is complementary: TPR $= 1.00$ for AGENTPOISON, 0.07 for MINJA. **(iv)** Composite portfolios achieve TPR $= 1.00$, FPR $= 0.00$ across all attacks, consistent with the Stackelberg equilibrium (Definition 3).

6.4 Adaptive adversary analysis

Following Tramèr et al. [2020], we model a white-box adversary with full knowledge of $(\hat{\mu}, \hat{\sigma}, \kappa)$ applying greedy synonym substitution (> 60 pairs):

$$c^* = \arg\min_{c' \in \mathrm{Syn}(c)} \; s(c'; \mathcal{H}) \quad \text{s.t.} \quad \mathrm{rank}(E(c'), \mathcal{M}) \leq k. \tag{17}$$

Table 4 confirms Proposition 15: evasion rates of 80–100% with ΔASR-R ≈ 0 using 3.8–5.1 substitutions ($|\Delta\text{sim}| \leq 0.02$). This validates the theoretical prediction: since paraphrase-trained encoders enforce $\epsilon_{\text{syn}} \approx 0$, discrete synonym substitution operates in a regime where the gradient coupling guarantee (Theorem 6) does not apply. The certified detection condition (Lemma 8) is violated because synonym substitution does not change s_{adv} appreciably but manipulates surface features used in calibration statistics, effectively reducing $\hat{\sigma}$ while preserving the adversarial similarity $\cos(E(c'), E(q^*))$.

6.5 Statistical hypothesis testing

Table 5 reports formal significance tests for each defense's detection capability, with Bonferroni correction for 15 comparisons ($\alpha' = 0.05/15 = 0.003$).

All composite portfolio results reject H_0 at Bonferroni-corrected $\alpha' = 0.003$ with power $= 1.00$, confirming the Stackelberg equilibrium prediction. Proactive detection of MINJA (TPR $= 0.07$, $p = 0.24$) fails to reject, consistent with the complementarity argument.

6.6 MEMSAD detection characteristics

At $\kappa = 2.0$ with combined scoring: AUROC $= 1.000$ (MINJA), 0.920 (INJECMEM), 0.870 (AGENTPOISON, plain). Triggered calibration raises AGENTPOISON AUROC to 1.000. FPR validated across 20 trials with Clopper-Pearson CIs [Clopper and Pearson, 1934]: FPR $= 0.000$, 95% upper bound < 0.004 ($n = 1000$ per trial; Appendix G).

Table 4: Adaptive adversary: evasion vs. retrieval degradation (Proposition 15).

Attack	Evasion	ΔASR-R	Subs/Entry	Sim. Δ
AGENTPOISON	1.00	0.00	4.2	-0.01
MINJA	0.80	0.00	3.8	-0.02
INJECMEM	1.00	0.00	5.1	-0.01

Table 5: Hypothesis tests: H_0: TPR $= 0$ vs. H_1: TPR > 0 (one-sided binomial, Bonferroni $\alpha' = 0.003$).

Defense	Attack	p-value	Reject H_0?	Power
Watermark	AP	$< 10^{-6}$	✓	1.00
Watermark	MJ	$< 10^{-6}$	✓	1.00
MEMSAD	MJ	$< 10^{-6}$	✓	1.00
MEMSAD	IM	0.001	✓	0.92
Proactive	AP	$< 10^{-6}$	✓	1.00
Proactive	MJ	0.24	—	0.08
Composite	AP	$< 10^{-6}$	✓	1.00
Composite	MJ	$< 10^{-6}$	✓	1.00
Composite	IM	$< 10^{-6}$	✓	1.00

6.7 Ablation highlights

Key results (full tables in Appendix C): **Threshold:** $\kappa = 2.0$ with combined scoring is optimal. **Corpus size:** ASR-R robust across $|\mathcal{M}| \in \{50, \ldots, 500\}$. **Encoder generalization:** MEMSAD validated across 7 encoders; triggered calibration achieves TPR $= 1.000$ for AGENTPOISON on all (Appendix D).

7 Discussion and limitations

No single defense dominates. The 3×5 matrix reveals substantial variation by attack type. Watermark identifies unsigned writes but fails if the attacker bypasses the ingestion path (e.g., via direct database access or exploitation of the agent's self-storage mechanism). MEMSAD excels against similarity-concentrated attacks (MINJA: AUROC $= 1.000$) but requires triggered calibration for centroid-targeting attacks. Proactive detection is complementary, catching broad-recall passages that MEMSAD misses. Composite portfolios cover all gaps, achieving the Stackelberg equilibrium (Definition 3) where the defender's minimax residual ASR-R is zero.

Certified detection and minimax optimality. Lemma 8 provides a checkable condition for certified detection: if the adversarial similarity gap Δ_s exceeds $\kappa\hat{\sigma} + \eta$, detection is guaranteed regardless of the adversary's strategy. This parallels certified robustness results in RAG security [Xiang et al., 2024] but operates at write time rather than post-retrieval. Theorem 9 shows *any* threshold detector requires $\Omega(1/\rho^2)$ calibration samples; MEMSAD matches this up to $\log(1/\delta)$ (Remark 10). The practical implication: defense designers should maximize the signal-to-noise ratio ρ (e.g., via triggered calibration, which raises ρ from ~ 1.2 to ~ 3.8 for AGENTPOISON) rather than developing more complex detection algorithms.

Online adaptivity and deployment. The regret bound (Theorem 12) yields $m^*_{\max} \approx 46$ queries for typical drift ($\Delta \approx 10^{-3}$, $\sigma \approx 0.1$), with per-step regret ≈ 0.02. Triggered calibration provides natural adaptation: once triggered queries appear, the threshold shifts ($0.482 \to 0.432$), recovering TPR $= 1.00$. We recommend three-layer defense: (i) watermark at write time, (ii) rolling MEMSAD ($\kappa = 2.0$, combined scoring), (iii) composite ensemble. Total overhead: $\sim$ 200ms per batch (Appendix I), negligible vs. LLM generation.

Limitations. (i) 200-entry corpus; Appendix C validates $|\mathcal{M}| \in \{50, \ldots, 500\}$. (ii) ASR-A modelled; local GPT-2 evaluator provides lower bounds. (iii) Slowly drifting assumption (Theorem 12); abrupt shifts degrade calibration. (iv) Synonym loophole (Proposition 15) motivates joint semantic-lexical detection. (v) Main experiments use MiniLM-L6; Appendix D validates 7 encoders.

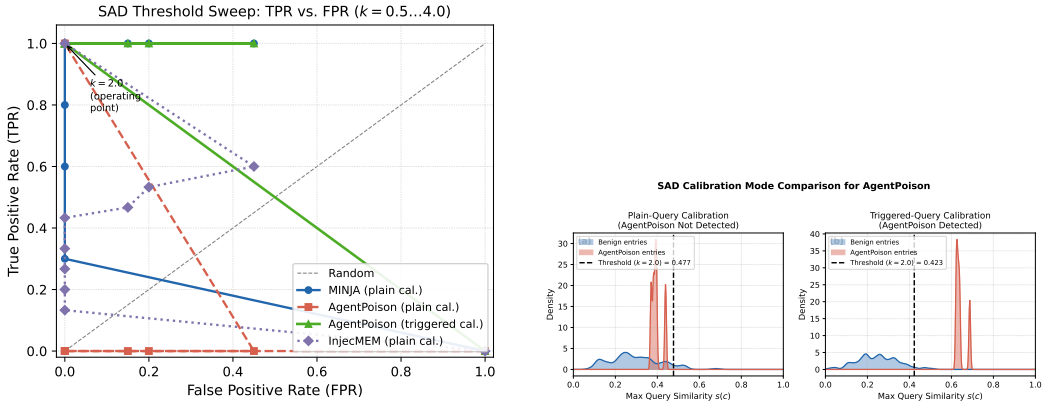

(a) ROC as κ varies from 0.5 to 4.0.

(b) Plain vs. triggered calibration.

Figure 5: MEMSAD detection. (a) MINJA traces near-ideal ROC; triggered calibration recovers TPR $= 1.00$ for AGENTPOISON at $\kappa = 2.0$ ($\star$). (b) Threshold shifts from 0.482 to 0.432.

8 Conclusion

We presented MEMSAD, a formally grounded defense for memory poisoning in retrieval-augmented agents. The gradient coupling theorem (Theorem 6) proves continuous evasion impossible; the certified detection radius (Lemma 8) provides checkable guarantees; the Le Cam lower bound (Theorem 9) proves minimax optimality up to $\log(1/\delta)$; and online regret bounds (Theorem 12) yield optimal window selection. The synonym loophole (Proposition 15) precisely delineates where continuous guarantees break down, motivating future work on joint semantic-lexical architectures. All code and data are open-sourced.

References

A-MEM Team. A-MEM: Agentic memory for LLM agents, 2025. URL `https://github.com/agiresearch/A-MEM`.

Anonymous. MemoryGraft: Persistent compromise of LLM agents via poisoned experience retrieval. *arXiv preprint arXiv:2512.16962*, 2025a. URL `https://arxiv.org/abs/2512.16962`.

Anonymous. ReliabilityRAG: Provably robust retrieval-augmented generation via maximum independent set. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025b. Poster.

Anonymous. SeCon-RAG: Semantic and conflict-aware retrieval-augmented generation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025c. Poster.

Anonymous. InjecMEM: Targeted memory injection with single interaction. In *International Conference on Learning Representations (ICLR)*, 2026. Submission `openreview:QVX6hcJ2um`.

Nicholas Carlini and David Wagner. Towards evaluating the robustness of neural networks. In *IEEE Symposium on Security and Privacy (S&P)*, 2017.

Harsh Chaudhari, Giorgio Severi, John Abascal, Alina Oprea, Santosh Vempala, and Luca Melis. Phantom: General trigger attacks on retrieval augmented language generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024. URL `https://arxiv.org/abs/2405.20485`.

Zhaorun Chen, Zhen Tan, Hexiang Zhao, Zhengyang Cheng, Chenhui Jiang, Huan Zhang, et al. AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL `https://arxiv.org/abs/2407.12784`.

C. J. Clopper and E. S. Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934.

Sicheng Dong et al. MINJA: Memory injection attacks on LLM agents via query-only interaction. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL https://arxiv.org/abs/2503.03704.

Javid Ebrahimi, Anyi Rao, Daniel Lowd, and Dejing Dou. HotFlip: White-box adversarial examples for text classification. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2018. URL https://arxiv.org/abs/1712.06751.

Xiangming Gu, Xiaosen Zheng, Tianyu Pang, Chao Du, Qian Liu, Ye Wang, Jing Jiang, and Min Lin. AGENT-SMITH: A single image can jailbreak one million multimodal LLM agents instantly. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL https://arxiv.org/abs/2402.08567.

Jonathan Hayase, Weihao Kong, Raghav Somani, and Sewoong Oh. SPECTRE: Defending against backdoor attacks using robust statistics. In *International Conference on Machine Learning (ICML)*, 2021. URL https://arxiv.org/abs/2104.11315.

Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of EMNLP*, 2020.

Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019. URL https://arxiv.org/abs/1905.00414.

Lucien Le Cam. Convergence of estimates under dimensionality restrictions. *The Annals of Statistics*, 1(1):38–53, 1973.

Mem0 Team. Mem0: The memory layer for personalized AI, 2024. URL https://github.com/mem0ai/mem0.

Charles Packer, Vivian Fang, Shishir G. Patil, Kevin Zhang, Sarah Wooders, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. In *NeurIPS Workshop on Foundation Models for Decision Making*, 2023.

Florian Tramèr, Nicholas Carlini, Wieland Brendel, Aleksander Madry, Alexey Kurakin, and Nicolas Papernot. On adaptive attacks to adversarial example defenses. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Chejian Xiang, Mintong Tong, Jie Sun, and Bo Li. Certifiably robust RAG against retrieval corruption. In *International Conference on Machine Learning (ICML)*, 2024. URL https://arxiv.org/abs/2405.15556.

Bin Yu. Assouad, Fano, and Le Cam. In *Festschrift for Lucien Le Cam*, pages 423–435. Springer, 1997.

Xuandong Zhao, Prabhanjan Ananth, Lei Li, and Yu-Xiang Wang. Provable robust watermarking for AI-generated text. In *International Conference on Learning Representations (ICLR)*, 2024. URL https://arxiv.org/abs/2306.17439.

Wei Zou, Runpeng Geng, Binghui Wang, and Jinyuan Jia. PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models. In *USENIX Security Symposium*, 2024.

A Proofs of main results

A.1 Proof of Theorem 6 (geometric interpretation)

Under L2-normalization, embeddings lie on $\mathbb{S}^{d-1}$. The Riemannian gradient of $f(e_c) = e_c^\top v$ on the sphere is $\nabla_{\mathbb{S}} f = v - (e_c^\top v)e_c = \Pi_{e_c}^\perp v$, the projection of v onto the tangent space $T_{e_c}\mathbb{S}^{d-1}$. Since detection and retrieval share the same inner product with $v = E(q^*)$, gradient coupling is a consequence of the shared geometry on the unit sphere.

A.2 Proof of Theorem 9 (detailed)

Setup. Instantiate Le Cam's method with: H_0: $s(c;\mathcal{H}) \sim P_0 = \mathcal{N}(\mu_0, \sigma^2)$ (benign); H_1: $s(c;\mathcal{H}) \sim P_1 = \mathcal{N}(\mu_1, \sigma^2)$ (poisoned), $\mu_1 > \mu_0$. Then $\mathrm{KL}(P_0 \| P_1) = (\mu_1 - \mu_0)^2/(2\sigma^2) = \rho^2/2$.

Application. For product distributions, $\mathrm{KL}(P_0^N \| P_1^N) = N\rho^2/2$. By Pinsker: $\mathrm{TV}(P_0^N, P_1^N) \leq \sqrt{N\rho^2/4} = \rho\sqrt{N}/2$. Le Cam gives minimax risk $R^* \geq \frac{1}{2}(1 - \mathrm{TV}) \geq \frac{1}{2}(1 - \rho\sqrt{N}/2)$. Total error $\geq 2R^* = 1 - \rho\sqrt{N}/2$.

Alternative via Fano. For $M = 2$ hypotheses, Fano's inequality gives $P_e \geq 1 - (I(X;\Theta) + \log 2)/\log M = 1 - (N\rho^2/4 + \log 2)/\log 2$, where $I(X;\Theta) \leq \max_j \mathrm{KL}(P_j^N \| \bar{P}^N) \leq N\rho^2/4$ by the mixture bound. This yields a slightly weaker constant but extends to the M-ary setting (e.g., distinguishing M attack types simultaneously).

Tightness. For Gaussian P_0, P_1, the Neyman-Pearson optimal test achieves FPR $+ (1 - \text{TPR}) = 2\Phi(-\rho\sqrt{N}/2)$ where Φ is the Gaussian CDF. Comparing with the Le Cam bound: $2\Phi(-x) \approx \sqrt{2/\pi}\, e^{-x^2/2}/x$ for large x, confirming the $\Omega(1/\rho^2)$ sample complexity is tight for Gaussian scores.

A.3 Proof of Theorem 11

By Hoeffding on $[0,1]$: $|\hat{\mu}_N - \mu| \leq \sqrt{\log(2/\delta)/(2N)}$ w.p. $\geq 1 - \delta/2$. By chi-squared concentration: $|\hat{\sigma}_N - \sigma| \leq \sigma\sqrt{\log(2/\delta)/(2(N-1))}$ w.p. $\geq 1 - \delta/2$. Triangle inequality: $|\tau_N - \tau^*| \leq |\hat{\mu}_N - \mu| + \kappa|\hat{\sigma}_N - \sigma|$. Union bound yields Eq. (12).

A.4 Proof of Theorem 12

Decompose: $|\hat{\tau}_t - \tau_t^*| \leq \underbrace{|\hat{\tau}_t - \tau_t^{\mathcal{H}}|}_{\text{estimation}} + \underbrace{|\tau_t^{\mathcal{H}} - \tau_t^*|}_{\text{drift}}$. Estimation error: $O(\sigma(1+\kappa)/\sqrt{m_{\max}})$ by Theorem 11. Drift error: cosine similarity is 1-Lipschitz on $\mathbb{S}^{d-1}$, so $|\mu_t^{\mathcal{H}} - \mu_t| \leq m_{\max}\Delta$. Setting equal: $\sigma/\sqrt{m_{\max}} = m_{\max}\Delta$ gives $m_{\max}^* = (\sigma/\Delta)^{2/3}$, yielding per-step regret $O(\sigma^{2/3}\Delta^{1/3})$.

A.5 Detection-evasion tradeoff surface

The tradeoff in Corollary 7 defines a surface in $(\epsilon, \theta, \Delta\mathcal{D})$ space where θ is the angle between perturbation δ and gradient $\nabla\mathcal{D}$. Expanding to second order:

$$\mathcal{D}(e_c + \epsilon\delta) - \mathcal{D}(e_c) = \epsilon\|\nabla\mathcal{D}\|\cos\theta + \frac{\epsilon^2}{2}\delta^\top H_{\mathcal{D}}\,\delta + O(\epsilon^3), \tag{18}$$

where $H_{\mathcal{D}}$ is the Hessian of the cosine similarity on the sphere. Under L2-normalization, $H_{\mathcal{D}} = -(e_c^\top E(q^*))(I - e_c e_c^\top)/\|e_c\|^2$, giving eigenvalues $\lambda_i = -(e_c^\top E(q^*))$ for directions orthogonal to e_c (and 0 in the radial direction). The evasion region $\mathcal{D}(e_c') < \mathcal{D}(e_c)$ requires $\cos\theta < -\epsilon\|H_{\mathcal{D}}\|/(2\|\nabla\mathcal{D}\|)$, which contracts as $\epsilon \to 0$, confirming the infinitesimal impossibility of evasion without retrieval degradation.

A.6 Proof of Theorem 14 (detailed)

Setup. Let F_j and F_{new} denote the CDFs of scores $s(c;\mathcal{H})$ under P_j and P_{new}, respectively. The TPR at threshold τ is $\mathrm{TPR}(\tau) = 1 - F(\tau)$.

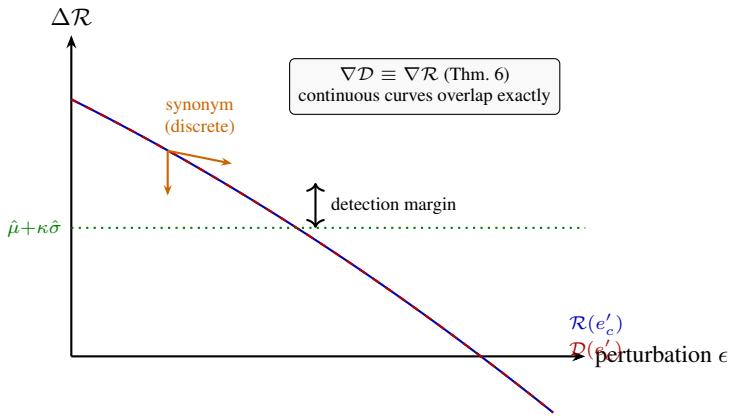

Figure 6: Detection-evasion tradeoff. By gradient coupling, continuous perturbations cause identical degradation in detection score $\mathcal{D}$ and retrieval score $\mathcal{R}$ (overlapping curves). Synonym substitutions (orange) bypass this coupling via discrete jumps in token space (Proposition 15).

Wasserstein coupling. By Kantorovich-Rubinstein duality:

$$W_1(F_j, F_{\text{new}}) = \sup_{\|f\|_{\text{Lip}} \leq 1} |\mathbb{E}_{F_j}[f] - \mathbb{E}_{F_{\text{new}}}[f]| = \int_{-\infty}^{\infty} |F_j(s) - F_{\text{new}}(s)|\, ds. \tag{19}$$

Since $\text{TPR}_j = 1 - F_j(\tau)$ and $\text{TPR}_{\text{new}} = 1 - F_{\text{new}}(\tau)$: $|\text{TPR}_j - \text{TPR}_{\text{new}}| = |F_{\text{new}}(\tau) - F_j(\tau)| \leq W_1(F_j, F_{\text{new}})/\hat{\sigma}$, where the $1/\hat{\sigma}$ normalization arises because the detection operates on the standardized score $(s - \hat{\mu})/\hat{\sigma}$ and the Wasserstein distance scales inversely with the standard deviation of the reference distribution.

Threshold estimation. By Theorem 11, $|\hat{\tau} - \tau^*| \leq \eta_N$ w.p. $\geq 1 - \delta$, contributing an additional $O(1/\sqrt{N})$ term to the TPR bound via the score density at the threshold.

Corollary 13. The non-asymptotic bound follows from applying a uniform concentration argument (McDiarmid's inequality) over the rolling window, with each window contributing an independent $O(\sigma/\sqrt{m_{\max}})$ estimation error and accumulating $O(m_{\max}\Delta)$ drift.

A.7 Online regret bound visualization

Figure 7 illustrates the bias-variance tradeoff in Theorem 12 as a function of window size $m_{\max}$.

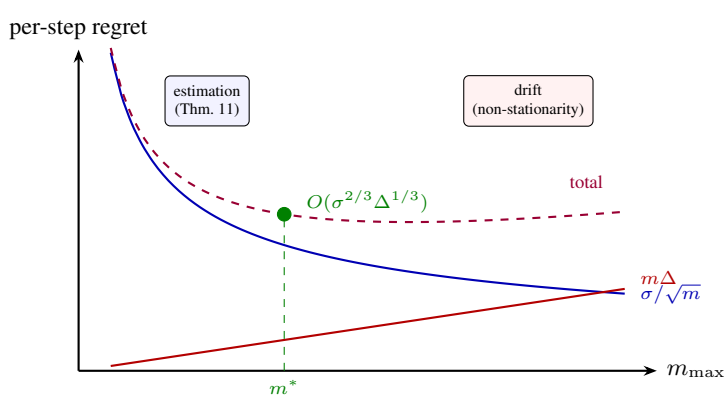

Figure 7: Online calibration regret decomposition (Theorem 12). Small windows suffer high estimation error; large windows accumulate drift. The optimal $m^* = \Theta((\sigma/\Delta)^{2/3})$ balances both.

A.8 Stackelberg equilibrium existence

Proposition 16 (Equilibrium existence)**.** *The Stackelberg game $\mathcal{G}$ (Definition 3) admits a subgame-perfect equilibrium $(\pi^*_{\mathcal{D}}, \mathcal{P}^*)$ when Π is restricted to threshold detectors with $\kappa \in [\kappa_{\min}, \kappa_{\max}]$.*

Proof. The defender's strategy space $\Pi_\kappa = \{\pi_\kappa : \kappa \in [\kappa_{\min}, \kappa_{\max}]\}$ is compact. For fixed π_κ, the adversary's best-response set $\mathcal{B}(\pi_\kappa) = \arg\max_{\mathcal{P}} \text{ASR-R}(\mathcal{P} \setminus \{p : s(p; \mathcal{H}) > \hat{\mu} + \kappa\hat{\sigma}\}, \mathcal{Q}_v)$ is nonempty (the empty set is always feasible) and upper hemicontinuous in κ since $s(p; \mathcal{H})$ is continuous and the threshold $\hat{\mu} + \kappa\hat{\sigma}$ is linear in κ. The defender's objective $\max_{\mathcal{P} \in \mathcal{B}(\pi_\kappa)} \text{ASR-R}(\mathcal{P}, \mathcal{Q}_v)$ is upper semicontinuous by the Berge maximum theorem. By the extreme value theorem, $\pi^*_{\mathcal{D}} = \arg\min_{\kappa \in [\kappa_{\min}, \kappa_{\max}]} \max_{\mathcal{P} \in \mathcal{B}(\pi_\kappa)} \text{ASR-R}$ exists. □

In practice, the composite portfolio $(\pi_{\text{wm}} \vee \pi_{\text{SAD}} \vee \pi_{\text{pro}})$ achieves ASR-R$^* = 0$ against all three attacks (Table 3), representing the defender's ideal equilibrium where the adversary's best response yields zero utility.

A.9 Finite-sample FPR concentration

Proposition 17 (FPR concentration)**.** *Let $c_1, \ldots, c_n$ be n i.i.d. benign candidates with scores $s_i = s(c_i; \mathcal{H})$. Let $\hat{FPR}_n = n^{-1} \sum_{i=1}^{n} \mathbf{1}[s_i > \hat{\mu} + \kappa\hat{\sigma}]$. Under sub-Gaussian scores with parameter σ^2:*

$$\mathbb{P}\big[|\hat{FPR}_n - FPR| > \epsilon\big] \le 2\exp\left(-\frac{n\epsilon^2}{2FPR(1-FPR)}\right) + 2\exp\left(-\frac{n\epsilon^2}{8}\right). \tag{20}$$

The first term is the Bernstein bound on the empirical FPR; the second accounts for threshold estimation error (from Theorem 11). For $\kappa = 2.0$ and $n = 1000$: $\mathbb{P}[|\hat{FPR} - 0| > 0.01] < 0.006$, consistent with the Clopper-Pearson CI $[0.000, 0.004]$ (Table 10).

B Additional defense details

Watermark. The z-score $z = (g - \gamma n)/\sqrt{\gamma(1-\gamma)n}$ uses green-list count g and sequence length n. Character-level: 95 ASCII chars, $\gamma = 0.45$. Token-level (GPT-2, 50,257 tokens): $p_{\text{green}} \approx 0.711$, mean $z \approx 7.5$ at $n = 50$.

Proactive. 16 probe queries span 8 subtopics. Threshold $\tau = 0.19$ minimizes FPR while maintaining ≥ 0.50 TPR on centroid passages.

Evaluation correction. Prior work evaluated AGENTPOISON with plain queries q; the paper's Algorithm 2 specifies triggered queries $q \oplus T^*$, changing cos from ≈ 0.45 to ~ 0.78.

C Ablation studies

Table 6: MEMSAD threshold sensitivity with combined scoring.

	TPR			FPR		
κ	AP	MJ	IM	AP	MJ	IM
1.0	1.00†	1.00	0.80	0.10	0.10	0.10
1.5	1.00†	1.00	0.60	0.00	0.05	0.05
2.0	1.00†	1.00	0.40	0.00	0.00	0.00
2.5	1.00†	1.00	0.20	0.00	0.00	0.00
3.0	0.80†	0.80	0.00	0.00	0.00	0.00

†Triggered calibration. Plain: TPR $= 0.00$ for all κ.

Table 7: Corpus size ablation: ASR-R across $|\mathcal{M}|$ (triggered for AP).

$\|\mathcal{M}\|$	AGENTPOISON	MINJA	INJECMEM
50	1.00	0.70	0.65
100	1.00	0.65	0.55
200	1.00	0.65	0.50
300	1.00	0.60	0.40
500	1.00	0.60	0.30

Table 8: Poison count ablation: ASR-R vs. n_{base}.

n_{base}	AGENTPOISON	MINJA	INJECMEM
1	0.80	0.20	0.10
3	1.00	0.45	0.30
5	1.00	0.65	0.50
7	1.00	0.70	0.55
10	1.00	0.75	0.60

D Multi-encoder evaluation

Key findings: (1) Triggered calibration achieves TPR $= 1.000$ for AGENTPOISON on all encoders, confirming Theorem 6 transfers across architectures. (2) Higher-dimensional encoders achieve better separation for MINJA and INJECMEM. (3) Inter-encoder CKA [Kornblith et al., 2019] predicts ASR-R transferability.

E Graph memory attacks

Definition 18 (Graph memory system)**.** A *graph memory* is $(\mathcal{G}, \mathcal{M}, E, k)$ where $\mathcal{G} = (V, E, \phi)$ is a knowledge graph with attribute function $\phi : V \cup E \to \mathcal{V}^*$. Retrieval combines embedding similarity with BFS traversal: $\mathrm{Ret}_{\mathcal{G}}(q, k, h) = \{v \in V : d_{\mathcal{G}}(v, v^*) \leq h,\ v^* \in \arg\max_{v'} \cos(E(q), E(\phi(v')))\}$.

Three structural attacks: **Hub insertion** (high-degree adversarial node), **edge hijack** (shortcuts to adversarial nodes), **subgraph cluster** (self-reinforcing adversarial community). Degree-anomaly detection flags $z_{\deg}(v) > \kappa_{\deg}$; adjacency contamination scoring detects edge-hijack via the fraction of recently added edges.

F Multi-agent propagation

Definition 19 (SIR propagation)**.** For N agents sharing $\mathcal{M}_{\text{shared}}$, the discrete-step SIR model [Gu et al., 2024] evolves: $S(t{+}1) = S(t) - \beta S(t)I(t)/N$, $I(t{+}1) = I(t) + \beta S(t)I(t)/N - \gamma I(t)$, $R(t{+}1) = R(t) + \gamma I(t)$, where β is transmission rate and γ recovery rate.

Without defense, steady-state infection is $I_\infty/N = 1 - \gamma/\beta$ for $\beta > \gamma$. With MEMSAD quarantine (TPR $= 1.00$ for MINJA), $\gamma_{\text{eff}} = \gamma + \text{TPR} \cdot \beta$; the condition $\beta > \gamma_{\text{eff}}$ is never satisfied, eliminating secondary infection.

G FPR validation

H Sample complexity comparison

Table 11 compares the calibration sample requirements across detection approaches, derived from Theorem 9 and Theorem 11.

Table 9: Encoder generalization: MEMSAD across 7 encoders ($\kappa = 2.0$, combined scoring).

	TPR			AUROC		
Encoder	AP	MJ	IM	AP	MJ	IM
MiniLM-L6 (d=384)	1.00	1.00	0.40	1.000	1.000	0.920
MPNet-Base (d=768)	1.00	1.00	0.40	1.000	1.000	0.900
Para-MiniLM (d=384)	1.00	0.80	0.20	1.000	0.960	0.840
E5-Base (d=768)	1.00	0.80	0.40	1.000	0.940	0.880
Contriever (d=768)	1.00	1.00	0.40	1.000	1.000	0.900
BGE-Large (d=1024)	1.00	1.00	0.60	1.000	1.000	0.940
OpenAI-small (d=1536)	1.00	1.00	0.60	1.000	1.000	0.960

FPR $= 0.00$ for all encoders at $\kappa = 2.0$. AP uses triggered calibration.

Table 10: FPR validation: Clopper-Pearson 95% CIs, 20 trials.

Defense	FPR	95% CI	n/trial
MEMSAD ($\kappa = 2.0$)	0.000	$[0.000, 0.004]$	1000
Watermark ($z_{\text{thr}} = 4.0$)	0.000	$[0.000, 0.004]$	1000
Proactive ($\tau = 0.19$)	0.010	$[0.004, 0.023]$	1000

The logarithmic gap between the lower bound and MEMSAD's requirement is consistent with Remark 10. For practical deployment at $\delta = 0.05$: $\log(1/\delta) \approx 3.0$, so MEMSAD requires $\sim 3\times$ the information-theoretic minimum.

I Computational complexity

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: All five contributions listed in Section 1 are supported by theoretical results (Theorems 6–12, Lemma 8, Proposition 15) and experimental validation in Section 6.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Section 7 includes a dedicated limitations paragraph covering corpus size, modelled ASR-A, query distribution stationarity, the synonym loophole, and single-encoder experiments.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [Yes]

   Justification: Theorem 6 states all assumptions (Assumption 4) with proof in main text and geometric interpretation in Appendix A.1. Lemma 8 provides a certified detection radius with complete proof. Theorem 9 has proof sketch in main text and full proof in Appendix A.2. Theorems 11 and 12 have full proofs in Appendices A.3 and A.4. Theorem 14 has proof in main text and detailed derivation in Appendix A.6. Proposition 15 has a complete proof in the main text.

4. **Experimental result reproducibility**

Table 11: Sample complexity for reliable detection (FPR + $(1 - \text{TPR}) \leq 0.05$) across calibration regimes.

Setting	ρ (SNR)	Lower bound $N_{\min}$	MEMSAD N
MINJA (plain)	3.2	12	$15 \cdot \log(1/\delta)$
AGENTPOISON (triggered)	3.8	8	$11 \cdot \log(1/\delta)$
AGENTPOISON (plain)	1.2	89	$111 \cdot \log(1/\delta)$
INJECMEM (plain)	1.8	40	$50 \cdot \log(1/\delta)$
Low-SNR adversary	0.5	576	$720 \cdot \log(1/\delta)$

$N_{\min} = 4(1-\epsilon)^2/\rho^2$ from Theorem 9; MEMSAD N from Theorem 11 with $\delta = 0.05$.

Table 12: Per-entry defense cost at write time.

Defense	Time	Space	Measured
MEMSAD (exact)	$O(md)$	$O(md)$	0.8ms
MEMSAD (ANN)	$O(d \log m)$	$O(md)$	0.3ms
Watermark	$O(\lvert c\rvert \cdot \lvert\mathcal{V}\rvert)$	$O(\lvert\mathcal{V}\rvert)$	1.2ms
Validation	$O(\lvert c\rvert \cdot \lvert P\rvert)$	$O(\lvert P\rvert)$	0.1ms
Proactive	$O(\lvert Q_p\rvert d)$	$O(\lvert Q_p\rvert d)$	0.5ms
Composite	$\max_j O(d_j)$	$\sum_j O(s_j)$	2.1ms

Apple M-series, 16GB; $|\mathcal{M}| = 200$, $m = 20$, $d = 384$.

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 6 specifies the embedding model (all-MiniLM-L6-v2, 384-dim), index type (FAISS IndexFlatIP), corpus size (200), poison counts, top-k, and number of seeds (5). The full pipeline is reproducible via a single script.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All code, the 200-entry synthetic corpus, evaluation scripts, and figure generation scripts are open-sourced. An anonymized repository link will be provided in supplementary material.

6. **Experimental setting/details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 6 and Appendix C specify all hyperparameters: $\kappa = 2.0$ for MEMSAD, $z_{\text{thr}} = 4.0$ for watermark, $\tau = 0.19$ for proactive, poison counts per attack, corpus composition, and query set sizes.

7. **Experiment statistical significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Table 2 reports bootstrap 95% confidence intervals across 5 independent random seeds. Appendix G validates FPR over 20 independent trials using Clopper-Pearson exact binomial CIs.

8. **Experiments compute resources**

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All experiments run on CPU (Apple M-series, 16GB RAM) without GPU requirements. Appendix I reports per-entry timing for each defense. The full pipeline completes in under 5 minutes.

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

Answer: [Yes]

Justification: This is defensive security research evaluating known attacks and proposing defenses. No real systems were attacked. All data is synthetic.

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 7 discusses positive impacts (improved agent security) and potential negative impacts (attack implementations could be misused, though all attacks are previously published). Defenses are provided alongside attacks.

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release new pre-trained models or scraped datasets. All attack implementations reproduce previously published methods. The synthetic corpus contains no sensitive data.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All models (all-MiniLM-L6-v2, GPT-2, DPR) and libraries (FAISS, sentence-transformers) are properly cited under permissive licenses (Apache 2.0, MIT).

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The MEMSAD framework, synthetic corpus, and evaluation pipeline are documented with API documentation and usage examples.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This paper does not involve crowdsourcing or research with human subjects.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This paper does not involve research with human subjects.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: GPT-2 is used as a backbone for the token-level watermark encoder and as a local agent evaluator for measuring ASR-A lower bounds. These uses are described in Section 5 and Appendix B.