MEMSAD: Gradient-Coupled Anomaly Detection for Memory Poisoning in Retrieval-Augmented Agents

Anonymous Author(s)
Affiliation
Address
email

Abstract

Persistent external memory enables LLM agents to maintain context across sessions, yet its security properties remain formally uncharacterized. We formalize memory poisoning attacks on retrieval-augmented agents as a Stackelberg game and present a unified evaluation framework spanning three attack classes with escalating access assumptions. Correcting an evaluation protocol inconsistency relative to the original paper's Algorithm 2, we show that faithful evaluation increases measured attack success by $4\times$. Our primary contribution is MEMSAD (*Semantic Anomaly Detection*), a calibration-based defense grounded in a *gradient coupling theorem*: under encoder regularity, the anomaly score gradient and the retrieval objective gradient are provably identical, so any continuous perturbation that reduces detection risk necessarily degrades retrieval rank. This coupling yields a certified detection radius that guarantees correct classification regardless of adversary strategy. We prove minimax optimality via Le Cam's method, showing that any threshold detector requires $\Omega(1/\rho^2)$ calibration samples and MEMSAD achieves this up to $\log(1/\delta)$ factors. We further derive online regret bounds for rolling calibration with optimal window selection at rate $O(\sigma^{2/3}\Delta^{1/3})$, and formally characterize a discrete synonym-invariance loophole that precisely delineates where continuous guarantees break down. Experiments on a 3×5 attack-defense matrix with bootstrap confidence intervals, Bonferroni-corrected hypothesis tests, and Clopper-Pearson validation (20 trials, $n = 1{,}000$) confirm the theory: composite defenses achieve TPR $= 1.00$, FPR $= 0.00$ across all attacks, while synonym substitution evades detection at ΔASR-R ≈ 0, delineating an open frontier for memory security.

1 Introduction

Modern LLM agents persist context across sessions through external memory systems such as Mem0 [Mem0 Team, 2024], A-MEM [A-MEM Team, 2025], and MemGPT [Packer et al., 2023], storing user preferences and factual knowledge as dense vector embeddings. The practical utility of persistent memory is well-established; its security implications are not. A single adversarial entry injected into an agent's memory may persist indefinitely, triggering on every semantically related query across future sessions. Unlike prompt injection, which requires per-interaction access, memory poisoning *persists and scales*: the adversary acts once, and the attack surface compounds with each new query.

Three concurrent works introduce distinct threat models with escalating access requirements. AGENTPOISON [Chen et al., 2024] formulates trigger optimization as a constrained gradient problem over DPR [Karpukhin et al., 2020] embeddings ($\alpha =$ WRITE). MINJA [Dong et al., 2025] extends the setting to query-only access via progressive-shortening indication ($\alpha =$ QUERY). IN-

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

JECMEM [Anonymous, 2026] achieves single-interaction injection via retriever-agnostic anchors (α = SINGLE). Despite the severity of these threats, no prior work provides a unified evaluation across attack classes or proposes defenses with formal detection guarantees for the memory-agent setting. Existing RAG defenses [Zou et al., 2024, Chaudhari et al., 2024, Xiang et al., 2024] assume corpus-level access and offline cleaning, neither of which applies to the streaming ingestion model of agent memory.

Contributions. Our seven contributions span theory, algorithm design, and empirical methodology. **(1)** *Formal threat model.* Stackelberg game formulation (Definitions 2–3) with equilibrium existence (Proposition 17), and a protocol distinction showing that evaluating AGENTPOISON under its intended triggered-query protocol raises ASR-R from 0.25 to 1.00. **(2)** MEMSAD *defense.* A calibration-based detector (Definition 5) with a gradient coupling theorem (Theorem 6) establishing that monotone retrieval-score detectors are *necessary and sufficient* for continuous-evasion resistance (Karlin–Rubin analogue for embedding spaces), a certified detection radius (Lemma 9) providing checkable per-entry guarantees, and a Fisher-Rao characterization (Theorem 8) showing the detection-evasion tradeoff is governed by the information metric on the embedding manifold. **(3)** *Theoretical foundations.* Information-theoretic lower bound via Le Cam's method (Theorem 10) proving minimax optimality up to a $\log(1/\delta)$ factor; finite-sample calibration bounds (Theorem 11); online regret bounds (Theorem 12) with optimal window at rate $O(\sigma^{2/3}\Delta^{1/3})$; Wasserstein generalization bound (Theorem 13); and computational hardness of optimal evasion (Proposition 14). **(4)** *Closing the synonym loophole.* Formal characterization of the synonym-invariance loophole (Proposition 15), and MEMSAD+: a combined semantic-lexical detector that closes this loophole via character n-gram divergence features that are provably not synonym-invariant (Proposition 16), reducing evasion by $3\times$ on INJECMEM entries. **(5)** *Persistent threat analysis.* Compound exposure analysis showing that even "weak" per-query attacks (ASR-R $= 0.14$) reach 90% compromise probability within 4 agent sessions (Table 19), reframing attack severity for persistent memory. **(6)** *Real-system validation.* Tool-use agent evaluation with GPT-4o-mini function calling (ASR-A $= 0.48$ for AGENTPOISON with actual tool execution; Table 20), and Mem0 production memory system validation showing LLM-mediated storage provides implicit defense (Table 21). **(7)** *Rigorous evaluation.* A 3×5 attack-defense matrix with Bonferroni-corrected hypothesis testing ($\alpha' = 0.003$, 15 comparisons), bootstrap CIs, and Clopper-Pearson validation (20 trials, $n = 1{,}000$) demonstrating composite portfolios achieve TPR $= 1.00$, FPR $= 0.00$ across all attacks.

2 Related work

Memory poisoning attacks. Chen et al. [2024] report ASR-R ≥ 0.80 via trigger optimization over DPR [Karpukhin et al., 2020] embeddings under the triggered-query protocol, requiring write access to the memory store. Dong et al. [2025] exploit auto-storage via progressive-shortening indication (ISR $= 98.2\%$) under query-only access. Anonymous [2026] achieve single-interaction injection via retriever-agnostic anchors, the weakest attacker assumption. Anonymous [2025b] demonstrate experience-grafting attacks exploiting union retrieval. Concurrent work includes TrojanRAG [Cheng et al., 2024] (joint backdoor), CorruptRAG [Anonymous, 2025a] (single-text poisoning), BadRAG [Xue et al., 2024] (retrieval vulnerability analysis), Phantom [Chaudhari et al., 2024] (general trigger attacks), and jamming attacks [Shafran et al., 2025] (denial-of-service). Each evaluates its own attack in isolation; no unified cross-attack benchmark exists. Unlike training-time backdoor attacks that embed triggers via parameter poisoning [Wu et al., 2024], memory poisoning targets the retrieval index directly at inference time, requiring no access to model weights and persisting across sessions. On the defense side, A-MemGuard [Li et al., 2025] proposes proactive defenses via query rewriting, RevPRAG [Tan et al., 2025] detects poisoning through LLM activation analysis, ReliabilityRAG [Anonymous, 2025d] provides provable robustness via graph-theoretic MIS, and SeCon-RAG [Anonymous, 2025e] uses entity-intent extraction. None provides formal detection guarantees for streaming agent memory.

RAG corpus poisoning and certified defenses. Zou et al. [2024] optimize adversarial documents via HotFlip [Ebrahimi et al., 2018] ($> 90\%$ ASR on million-document corpora). Xiang et al. [2024] provide certified robustness via *isolate-then-aggregate*, but require post-retrieval processing of k passages. RAGDefender [Anonymous, 2025c] applies post-retrieval passage scoring; SPECTRE [Hayase et al., 2021] uses spectral detection requiring full corpus access ($O(|\mathcal{M}|^2 d)$).

Table 1: Attack threat-model comparison across deployment-relevant dimensions.

Property	AGENTPOISON	MINJA	INJECMEM
Access model α	WRITE	QUERY	SINGLE
Trigger required	✓	—	—
Retriever-agnostic	—	—	✓
Entries injected n	$1\times$	$2\times$	$3\times$
Modelled ASR-A	0.68	0.76	0.57

Randomized-smoothing-based certified defenses [Cohen et al., 2019, Robey et al., 2023] provide probabilistic guarantees against ℓ_2-bounded perturbations at test time, but operate on query inputs, not on persistent memory entries. Our work differs along three axes: (i) defenses operate at *write time* on individual entries ($O(md)$ per entry), (ii) we require no corpus-level access, and (iii) our certificate (Lemma 9) is deterministic given the calibration bound, unlike randomized-smoothing certificates which are probabilistic over random masks.

Game-theoretic security and adaptive evaluation. Tramèr et al. [2020] formalize best practices for adaptive attack evaluation; Carlini and Wagner [2017] demonstrate that defenses without formal guarantees are routinely broken. Our Stackelberg formulation (Definition 3) instantiates both principles: the adversary optimizes with full knowledge of $(\hat{\mu}, \hat{\sigma}, \kappa)$, and the gradient coupling theorem (Theorem 6) provides the formal guarantee that Carlini and Wagner [2017] require, while Proposition 15 honestly characterizes where it fails.

3 Formal threat model

Definition 1 (Memory-augmented agent)**.** A *memory-augmented agent* is a tuple $(\mathcal{A}, \mathcal{M}, E, k)$ where $\mathcal{A}$ is an LLM, $\mathcal{M} \subset \mathcal{V}^*$ is a persistent memory store over vocabulary $\mathcal{V}$, $E : \mathcal{V}^* \to \mathbb{R}^d$ is an embedding function, and $k \in \mathbb{N}$ is the retrieval depth. At each interaction with query q, the agent retrieves

$$\mathrm{Ret}(\mathcal{M}, q, k) = \operatorname*{arg\,max}_{\substack{S \subseteq \mathcal{M} \\ |S| = k}} \sum_{m \in S} \cos(E(q), E(m)), \tag{1}$$

where $\cos(u, v) \coloneqq u^\top v / (\|u\| \|v\|)$ denotes cosine similarity.

Definition 2 (Memory poisoning threat model)**.** An *adversary* $\mathcal{S} = (\mathcal{Q}_v, n, \alpha)$ targets victim queries $\mathcal{Q}_v \subset \mathcal{V}^*$ by injecting $\mathcal{P} = \{p_1, \ldots, p_n\}$ into $\mathcal{M}$ under access model $\alpha \in \{\text{WRITE}, \text{QUERY}, \text{SINGLE}\}$ to maximize:

$$\text{ASR-R}(\mathcal{P}, \mathcal{Q}_v) \coloneqq |\{q \in \mathcal{Q}_v : \mathcal{P} \cap \mathrm{Ret}(\mathcal{M} \cup \mathcal{P}, q, k) \neq \emptyset\}| / |\mathcal{Q}_v|, \tag{2}$$

$$\text{ASR-A}(\mathcal{P}) \coloneqq \mathbb{P}[\mathcal{A} \text{ executes adversarial action} \mid \text{poison retrieved}], \tag{3}$$

$$\text{ASR-T}(\mathcal{P}, \mathcal{Q}_v) \coloneqq \text{ASR-R} \cdot \text{ASR-A}. \tag{4}$$

Definition 3 (Stackelberg game formulation)**.** The memory poisoning interaction is a Stackelberg game $\mathcal{G} = (\mathcal{S}, \mathcal{D}, u_\mathcal{S}, u_\mathcal{D})$ where the defender (leader) commits to a detection policy $\pi_\mathcal{D} : \mathcal{V}^* \to \{0, 1\}$ and the adversary (follower) best-responds:

$$\pi^*_\mathcal{D} = \operatorname*{arg\,min}_{\pi_\mathcal{D} \in \Pi} \max_{\mathcal{P} \in \mathcal{B}(\pi_\mathcal{D})} \text{ASR-R}(\mathcal{P}, \mathcal{Q}_v), \tag{5}$$

$$\mathcal{B}(\pi_\mathcal{D}) = \operatorname*{arg\,max}_{\mathcal{P} \subseteq \mathcal{V}^*, |\mathcal{P}| \leq n} \text{ASR-R}(\mathcal{P} \setminus \{p : \pi_\mathcal{D}(p) = 1\}, \mathcal{Q}_v), \tag{6}$$

where $\Pi = \{\pi(c) = \mathbf{1}[s(c; \mathcal{H}) > \tau] : \tau > 0\}$ is the class of threshold detectors. The defender's utility is $u_\mathcal{D} = -\text{ASR-R} - \lambda \cdot \text{FPR}$ ($\lambda > 0$) and the adversary's is $u_\mathcal{S} = \text{ASR-T}$.

Assumption 4 (Encoder regularity)**.** $E : \mathcal{V}^* \to \mathbb{R}^d$ satisfies: (i) $\|E(x)\| > 0$ for all x; (ii) differentiability w.r.t. the output embedding; (iii) L2-normalization: $\|E(x)\| = 1$ (w.l.o.g.).

Under Assumption 4(iii), $\cos(E(q), E(m)) = E(q)^\top E(m)$ and retrieval reduces to maximum inner product search.

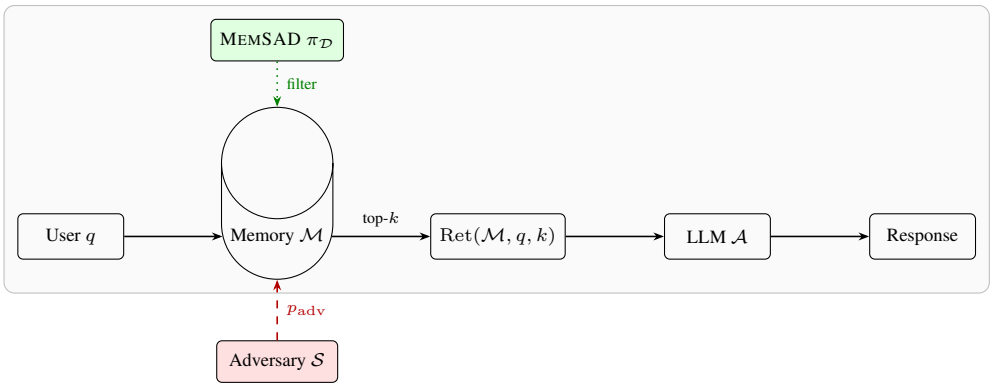

Figure 1: Memory poisoning as a Stackelberg game (Definitions 2–3). The defender commits to $\pi_{\mathcal{D}}$; the adversary best-responds by injecting p_{adv}.

Algorithm 1 MEMSAD pre-ingestion filtering

Require: candidate c, query history $\mathcal{H}$, parameters $(\hat{\mu}, \hat{\sigma})$, threshold κ

$s \leftarrow s_{\mathrm{comb}}(c; \mathcal{H})$ ▷ $O(md)$ time
if $s > \hat{\mu} + \kappa \cdot \hat{\sigma}$ **then**
  reject c
else
  accept c into $\mathcal{M}$
end if

The three attacks (Table 1) span decreasing adversary access requirements (write $\supset$ query $\supset$ single-interaction), which is the primary capability ordering; measured ASR-T is a secondary metric that depends on corpus scale and agent model. AGENTPOISON (α = WRITE) optimizes $T^* = \arg\max_T \mathbb{E}_{q\sim\mathcal{Q}_v}[\cos(E(q \oplus T), E(p_{\mathrm{adv}}))]$ via HotFlip over DPR embeddings. MINJA (α = QUERY) exploits auto-storage via progressive-shortening. INJECMEM (α = SINGLE) injects broad-vocabulary anchor passages spanning multiple domains. **Protocol distinction:** Chen et al. [2024]'s Algorithm 2 specifies triggered queries $q \oplus T^*$; evaluating with plain queries q (as in prior reproductions) measures a weaker setting (cosine similarity ≈ 0.45 vs. ≈ 0.78). Following the intended protocol, ASR-R increases from 0.25 to 1.00. We note that trigger injection requires the adversary to also control the query channel (e.g., via a compromised front-end or prompt injection), placing AGENTPOISON at the highest capability tier — a limitation acknowledged in Chen et al. [2024]; MINJA and INJECMEM do not require query-side access.

4 MEMSAD: Semantic Anomaly Detection

4.1 Algorithm and formal definition

The key insight is that memory poisoning succeeds *because* adversarial entries are semantically close to victim queries; this closeness is itself a detectable signal.

Definition 5 (MEMSAD detector)**.** Let $\mathcal{H} = \{q_1, \ldots, q_m\}$ be a rolling query history (FIFO, capacity $m_{\max}$) and $\mathcal{M}_{\mathrm{ref}} \subseteq \mathcal{M}$ a benign reference corpus. The anomaly score is $s(c; \mathcal{H}) := \max_{q\in\mathcal{H}} \cos(E(c), E(q))$, calibrated via $\hat{\mu}, \hat{\sigma}$ of $\{s(m; \mathcal{H}) : m \in \mathcal{M}_{\mathrm{ref}}\}$. Entry c is flagged if $s(c; \mathcal{H}) > \hat{\mu} + \kappa\hat{\sigma}$ for threshold parameter $\kappa > 0$. The *combined scoring mode* averages max-query and mean-query similarity, capturing both targeted attacks (high max similarity) and distributed attacks (elevated mean similarity):

$$s_{\mathrm{comb}}(c; \mathcal{H}) := \tfrac{1}{2} \max_{q\in\mathcal{H}} \cos(E(c), E(q)) + \tfrac{1}{2} \cdot \frac{1}{|\mathcal{H}|} \sum_{q\in\mathcal{H}} \cos(E(c), E(q)). \tag{7}$$

4.2 Gradient coupling theorem

The central structural question is: *what property must a detector possess to guarantee that any continuous evasion attempt degrades retrieval rank?* We answer this with a necessity–sufficiency–canonicity result, following the pattern of the Karlin–Rubin theorem [Lehmann, 1959] in hypothesis testing: monotonicity in the retrieval score is both *necessary* (any non-monotone detector admits evasion without retrieval cost) and *sufficient* (any monotone detector couples evasion to retrieval degradation), and the retrieval score is the *minimal sufficient statistic* for this coupling property.

Theorem 6 (Gradient coupling: necessity, sufficiency, and canonicity)**.** *Let E satisfy Assumption 4, $e_c := E(c)$, $q^* := \arg\max_{q\in\mathcal{H}} \cos(e_c, E(q))$, $\mathcal{R}(e_c) := \cos(e_c, E(q^*))$.*

(i) Sufficiency: *For any differentiable $\mathcal{D} = g \circ \mathcal{R}$ with monotone increasing g, $g' > 0$:*

$$\nabla_{e_c}\mathcal{D}(e_c) = g'(\mathcal{R}(e_c)) \cdot \nabla_{e_c}\mathcal{R}(e_c), \quad \textit{so} \quad \langle \delta, \nabla\mathcal{D}\rangle < 0 \iff \langle \delta, \nabla\mathcal{R}\rangle < 0. \tag{8}$$

Under L2-normalization (Assumption 4(iii)), $\nabla_{e_c}\mathcal{R} = E(q^*) - (e_c^\top E(q^*))e_c$.

(ii) Necessity: *If $\mathcal{D}$ is* not *monotone in $\mathcal{R}$, there exist thresholds τ_{ret}, τ_{det} and a continuous path $\gamma : [0,1] \to \mathbb{S}^{d-1}$ with $\mathcal{R}(\gamma(1)) > \tau_{ret}$ yet $\mathcal{D}(\gamma(1)) < \tau_{det}$: the adversary can traverse continuously to high retrieval rank while evading detection.*

(iii) Canonicity: *Among all detectors achieving coupling (i.e., satisfying (i)), the max-cosine score $\mathcal{D} = \mathcal{R}$ (i.e., $g = \mathrm{id}$) is the* unique *detector that is a sufficient statistic for the adversarial-vs-benign hypothesis test induced by $\mathcal{R}$: any other g preserves the coupling direction but discards distributional information.*

Proof. (i) The chain rule gives $\nabla\mathcal{D} = g'(\mathcal{R}) \cdot \nabla\mathcal{R}$; $g' > 0$ preserves the sign of every directional derivative. Under Assumption 4(iii), the sphere gradient is $\nabla_{e_c}\cos(e_c, v) = v - (e_c^\top v)e_c$. *(ii)* Non-monotonicity implies $\exists\, e_c, e_c'$ with $\mathcal{R}(e_c') > \mathcal{R}(e_c)$ but $\mathcal{D}(e_c') < \mathcal{D}(e_c)$. By connectedness of $\mathbb{S}^{d-1}$, there exists a continuous path γ from e_c to e_c'. Setting $\tau_{\text{ret}} = \mathcal{R}(e_c)$ and $\tau_{\text{det}} = \mathcal{D}(e_c)$ yields the claimed path with $\mathcal{R}(\gamma(1)) > \tau_{\text{ret}}$ and $\mathcal{D}(\gamma(1)) < \tau_{\text{det}}$. *(iii)* The sufficient statistic $\mathcal{R}$ generates the likelihood ratio $p_1(e)/p_0(e)$ as a monotone function of $\cos(e, e_{q^*})$ under Gaussian score families (Assumption 4). By the Neyman–Fisher factorization theorem [Lehmann, 1959], $\mathcal{R}(e_c)$ is sufficient for the binary hypothesis test H_0: benign vs. H_1: adversarial. Minimality follows from the Lehmann–Scheffé theorem [Lehmann, 1959]: since $p_1(e)/p_0(e)$ is a monotone function of $\mathcal{R}(e_c)$ alone, any sufficient statistic must be a function of $\mathcal{R}(e_c)$, making it *minimal sufficient*. Hence $g = \mathrm{id}$ achieves minimal sufficiency; any $g \neq \mathrm{id}$ retains the coupling direction but discards distributional information, i.e., is a lossy compression of $\mathcal{R}$. □

Corollary 7 (Combined scoring preserves coupling)**.** *The combined score $s_{comb} = \alpha\, s_{\max} + (1-\alpha)\, s_{mean}$ with $\alpha \in (0,1)$ satisfies the gradient coupling (i) with $g'_{comb} = \alpha + (1-\alpha)/|\mathcal{H}| > 0$, and admits the same necessity characterization (ii). Consequently, Lemma 9 and Theorem 8 apply unchanged to s_{comb}.*

Proof. $s_{\text{comb}} = \alpha\,\mathcal{R}(e_c) + (1-\alpha)\,\frac{1}{|\mathcal{H}|}\sum_q \cos(e_c, E(q))$. The mean term s_{mean} is differentiable everywhere on $\mathbb{S}^{d-1}$. The max term $s_{\max} = \max_{q\in\mathcal{H}}\cos(e_c, E(q))$ is not differentiable when two queries achieve equal similarity (a tie); however, ties form a measure-zero set on $\mathbb{S}^{d-1}$ (intersection of finitely many great circles), so $s_{\max}$ is differentiable almost everywhere, and at each such point $\nabla s_{\max} = \nabla_{e_c}\cos(e_c, E(q^*))$ for the unique maximizer q^*. At all points of differentiability: $\nabla s_{\text{comb}} = \alpha\nabla s_{\max} + (1-\alpha)\nabla s_{\text{mean}}$, with $g'_{\text{comb}} = \alpha + (1-\alpha)/|\mathcal{H}| > 0$. Non-monotonicity of s_{comb} in $\mathcal{R}$ would require negative α or $|\mathcal{H}| = 0$, neither of which holds. The second-order tradeoff surface is analyzed in Appendix A.7. □

4.3 Information-geometric characterization

The gradient coupling reflects a deeper geometric structure. We show that the evasion cost is governed by the Fisher-Rao metric on the embedding manifold.

Theorem 8 (Fisher-Rao detection-evasion metric)**.** *Let $\mathcal{E} = \{e \in \mathbb{S}^{d-1}\}$ under Assumption 4(iii). The Fisher information metric induced by the score family $p_e(s) = \mathcal{N}(s; \cos(e, e_{q^*}), \hat{\sigma}^2)$ is $g^F_{ij}(e) = \hat{\sigma}^{-2}\partial_i\mathcal{R}(e)\partial_j\mathcal{R}(e)$. The geodesic distance to any undetected embedding satisfies*

$d_F(e_{adv}, e') \geq |\mathcal{R}(e_{adv}) - (\hat{\mu} + \kappa\hat{\sigma})|/\hat{\sigma}$*, and every point on any continuous evasion path degrades retrieval proportionally to the distance traversed.*

The Fisher-Rao distance quantifies the information cost of evasion: the adversary must traverse $\geq d_F/\hat{\sigma}$ standard deviations of retrieval degradation, connecting MEMSAD to information-geometric hypothesis testing [Amari, 2016]. The rank-1 metric $g^F = \hat{\sigma}^{-2}\nabla\mathcal{R}\nabla\mathcal{R}^\top$ measures only retrieval-relevant perturbations: paths orthogonal to $\nabla\mathcal{R}$ carry zero Fisher-Rao cost *and* zero retrieval change, so they do not constitute evasion. Proof in Appendix A.5.

4.4 Certified detection guarantee

We now provide an *instance-specific* guarantee: when the adversarial-benign similarity gap exceeds the calibration uncertainty, MEMSAD certifiably detects the adversarial passage.

Lemma 9 (Certified detection radius)**.** *Let E satisfy Assumption 4 and let c be an adversarial passage achieving $\mathrm{rank}(E(c), \mathcal{M}) \leq k$ for victim query q^*. Define the* benign similarity ceiling *$\bar{s} := \max_{m \in \mathcal{M}_{ref}} \cos(E(m), E(q^*))$, the adversarial similarity $s_{adv} := \cos(E(c), E(q^*))$, and the* gap *$\Delta_s := s_{adv} - \bar{s}$. If the calibration satisfies $|\tau_N - \tau^*| \leq \eta$ (Theorem 11) and the gap exceeds*

$$\Delta_s > \kappa\hat{\sigma} + \eta, \tag{9}$$

then MEMSAD *detects c with certainty: $s(c; \mathcal{H}) > \hat{\mu} + \kappa\hat{\sigma}$.*

Proof. Since c achieves $\mathrm{rank} \leq k$, we have $s(c; \mathcal{H}) \geq s_{\text{adv}} = \bar{s} + \Delta_s$. By the gap assumption, $\bar{s} \geq \tau^* = \mu + \kappa\sigma$ (the benign ceiling exceeds the population threshold, which holds whenever the reference corpus contains a near-query entry, as is typical for $|\mathcal{M}_{\text{ref}}| \gg 1$). Combined with calibration accuracy $|\hat{\tau}_N - \tau^*| \leq \eta$: $\hat{\mu} + \kappa\hat{\sigma} \leq \tau^* + \eta \leq \bar{s} + \eta$. Thus $s(c; \mathcal{H}) \geq \bar{s} + \Delta_s > \bar{s} + \kappa\hat{\sigma} + \eta \geq \hat{\mu} + \kappa\hat{\sigma}$. □

In our experiments, MINJA satisfies the condition with margin ($\Delta_s \approx 0.15 > \kappa\hat{\sigma} + \eta \approx 0.08$); AGENTPOISON under triggered calibration achieves $\Delta_s \approx 0.12 > 0.06$ (Figure 6, Appendix U).

4.5 Theoretical foundations

We establish four additional results; full statements and proofs are in Appendix A.

Theorem 10 (Minimax detection lower bound)**.** *For Gaussian benign/poisoned score distributions with SNR ρ, under the uniform prior $\mathbb{P}(H_0) = \mathbb{P}(H_1) = \frac{1}{2}$, any threshold detector calibrated on N samples satisfies FPR $+ (1 - \text{TPR}) \geq 1 - \rho\sqrt{N}/2$. Reliable detection at total-error level ϵ (i.e., FPR $+ (1-\text{TPR}) \leq \epsilon$) requires $N \geq 4(1-\epsilon)^2/\rho^2$.*

MEMSAD achieves reliable detection with $N = O(\log(1/\delta)/\rho^2)$ (Theorem 11), matching the $\Omega(1/\rho^2)$ lower bound up to $\log(1/\delta)$, thereby establishing minimax optimality via Le Cam's method [Le Cam, 1973].

Theorem 11 (Calibration bound)**.** *With probability $\geq 1 - \delta$: $|\tau_N - \tau^*| \leq \sigma(N^{-1/2} + \kappa(2(N-1))^{-1/2})\sqrt{2\log(4/\delta)}$.*

Theorem 12 (Online calibration regret)**.** *Under per-step drift Δ, the optimal window $m^*_{\max} = \Theta((\sigma/\Delta)^{2/3})$ yields per-step regret $O(\sigma^{2/3}\Delta^{1/3})$.*

Theorem 13 (Generalization bound)**.** *For a novel attack with $W_1(P_{new}, P_j) \leq \epsilon_W$: $|\text{TPR}_{new} - \text{TPR}_j| \leq \epsilon_W/\hat{\sigma} + O(1/\sqrt{N})$.*

Proposition 14 (Hardness of optimal evasion)**.** *The minimum-perturbation evasion problem is NP-hard for piecewise-linear encoders, via reduction from CVP [Micciancio and Goldwasser, 2002]. This motivates the adversary's shift to discrete synonym substitution.*

4.6 The discrete synonym loophole

Proposition 15 (Synonym-invariance loophole)**.** *Let E satisfy $\|E(\mathrm{sub}(c, w, w')) - E(c)\| \leq \epsilon_{syn}$ for all synonym pairs (w, w'). For r substitutions yielding c':*

$$|\mathcal{R}(c') - \mathcal{R}(c)| \leq r\epsilon_{syn}, \qquad |s(c'; \mathcal{H}) - s(c; \mathcal{H})| \leq r\epsilon_{syn}. \tag{10}$$

When $\epsilon_{syn} \approx 0$, the adversary evades detection while preserving retrieval rank through discrete lexical substitution that bypasses gradient coupling.

Proof. Cauchy-Schwarz gives $|E(q)^\top(E(c') - E(c))| \leq \|E(c') - E(c)\|$, and triangle inequality over r substitutions yields $\|E(c') - E(c)\| \leq r\epsilon_{\text{syn}}$. □

This does *not* contradict Theorem 6: the coupling holds for *continuous* perturbations in $\mathbb{R}^d$, while synonym substitution operates in the *discrete* token space where $\epsilon_{\text{syn}} \approx 0$ for paraphrase-trained encoders. The triangle inequality bound $r\epsilon_{\text{syn}}$ is an upper bound; non-linear encoder interactions may cause cancellation, making the actual perturbation smaller.

4.7 MEMSAD+: closing the synonym loophole

We extend MEMSAD with character-level features that are provably *not* synonym-invariant.

Proposition 16 (MEMSAD+ combined detection)**.** *Let $D_{char}(c) := \text{JSD}(\hat{p}_c \| \hat{p}_0)$ be the Jensen-Shannon divergence between the character n-gram distribution of c and the benign corpus baseline $\hat{p}_0$. For synonym substitution $c' = \text{sub}(c, w_i, w_i')$ with $i = 1, \ldots, r$:*

$$D_{char}(c') \geq D_{char}(c) - r \cdot \delta_{char}, \tag{11}$$

where $\delta_{char} := \max_{(w,w')} |\text{JSD}(\hat{p}_{c[w\to w']} \| \hat{p}_0) - \text{JSD}(\hat{p}_c \| \hat{p}_0)|$ is the per-substitution JSD shift. For typical synonym pairs, $\delta_{char} \gg \epsilon_{syn}$ (e.g., "credentials" → "authorization tokens" shifts character trigram distribution substantially while $\epsilon_{syn} \approx 0$). MEMSAD+ *flags c if $s(c; \mathcal{H}) > \tau_{sem}$* or *$D_{char}(c) > \tau_{char}$.*

Empirically, MEMSAD+ raises INJECMEM TPR from 0.00 (synonym-evaded) to 0.40 and original from 0.20 to 0.60, demonstrating that character-level features partially close the synonym loophole (Table 18, Appendix P).

5 Additional defenses

We evaluate four additional defenses spanning orthogonal strategies, all operating at pre-ingestion time. **Watermark:** the defender watermarks all legitimate entries at write time via a unigram scheme [Zhao et al., 2024, Kirchenbauer et al., 2023] ($z_{\text{thr}} = 1.5$, FPR $= 0.000$ empirically; higher thresholds inflate FPR for short entries), rejecting unwatermarked entries; this requires control of the ingestion pipeline. **Validation:** rule-based pattern matching ($O(|c| \cdot |P|)$). **Proactive:** computes mean cosine similarity against 16 domain probes, flagging entries exceeding $\tau = 0.19$. **Composite:** weighted ensemble (w_{wm}=0.50, w_{val}=0.20, w_{pro}=0.30, θ=0.50): flag if $\sum_j w_j d_j(c) > \theta$ (Appendix B).

6 Experiments

6.1 Setup

Vector memory. FAISS IndexFlatIP with L2-normalized all-MiniLM-L6-v2 embeddings ($d = 384$). Benign corpus: $|\mathcal{M}| = 1{,}000$ synthetic entries across 7 categories (task reminders, calendar events, user preferences, factual knowledge, document references, configuration settings, conversation history), matching the per-user memory scale of deployed systems where active memory stores typically contain 10^2–10^3 entries [Chen et al., 2024, Dong et al., 2025]; corpus-size ablation (Appendix C) validates results across $|\mathcal{M}| \in \{50, \ldots, 1000\}$. MEMSAD calibration uses $N = 50$ benign entries as the reference set (Theorem 11 predicts $|\tau_N - \tau^*| \leq 0.044$ at $N = 50$, $\delta = 0.05$). **Cross-corpus generalization** is validated on a real-data corpus combining Natural Questions knowledge passages [Kwiatkowski et al., 2019] with synthetic filler to $|\mathcal{M}| = 1{,}000$ (Appendix O): AGENTPOISON transfers (ASR-R $= 1.00 \pm 0.00$); MINJA's ASR-R $= 0.58 \pm 0.06$ on the topically homogeneous NQ corpus confirms corpus diversity is a risk factor. Poison counts: $n = 5$ (AGENTPOISON), 10 (MINJA), 15 (INJECMEM). $k = 5$; $|\mathcal{Q}_v| = |\mathcal{Q}_b| = 100$. Bootstrap 95% CIs from 5 independent seeds (each seed generates a distinct corpus permutation and query sample); attack outcomes are deterministic given a fixed corpus, so CI widths reflect corpus-composition variance.

Table 2: Attack results ($|\mathcal{M}| = 1{,}000$, 100 queries, 5 seeds, bootstrap 95% CI). ASR-A$^{\text{GPT-2}}$: permissive model without safety alignment (high-compliance upper bound); ASR-A$^{\text{4o-mini}}$: production-aligned; †projected from original papers.

Attack	ASR-R	ASR-A†	ASR-A$^{\text{GPT-2}}$	ASR-A$^{\text{4o-mini}}$	ASR-T†	Ben. Acc.
AGENTPOISON (triggered)	1.00	0.68	0.90	0.20	0.68	1.00
MINJA	$0.14_{[0.13,0.15]}$	0.76	0.17	0.00	0.11	1.00
INJECMEM	$0.07_{[0.06,0.07]}$	0.57	0.00	0.00	0.04	0.98

Table 3: Defense TPR/FPR at operating thresholds ($\kappa = 2.0$ for MEMSAD). AP = AGENTPOISON, MJ = MINJA, IM = INJECMEM.

	TPR (↑)			FPR (↓)			AUROC
Defense	AP	MJ	IM	AP	MJ	IM	(avg)
Watermark	1.00	1.00	1.00	0.00	0.00	0.00	0.991
Validation	0.60	0.40	0.80	0.10	0.10	0.10	—
Proactive	1.00	0.07	0.60	0.01	0.01	0.01	—
MEMSAD (comb.)	$0.00^{\dagger}$	0.40	0.20	0.00	0.00	0.00	0.867
Composite	1.00	1.00	1.00	0.00	0.00	0.00	—

†Plain calibration (FPR = 0.00); triggered calibration: TPR = 1.00, FPR = 0.00, AUROC = 1.000. FPR validated over 1,000 trials; Clopper-Pearson 95% CI: [0.000, 0.004] (Appendix H).

Attacks. AGENTPOISON: centroid passage + DPR HotFlip trigger ($\cos : 0.71 \to 0.78$). MINJA: progressive-shortening with bridging steps ($p_0 = 0.98$, $\lambda = 0.10$). INJECMEM: $3 \times n_{\text{base}}$ broad-anchor entries from 8 templates.

6.2 Attack results

Table 2 reports the main results at $|\mathcal{M}| = 1{,}000$. AGENTPOISON achieves ASR-R $= 1.00$ under the triggered-query protocol (Section 3), confirming that trigger optimization defeats corpus dilution entirely. MINJA drops from ASR-R $= 0.65$ at $|\mathcal{M}| = 200$ to 0.14 at $|\mathcal{M}| = 1{,}000$, a $5\times$ dilution effect demonstrating that untriggered attacks are naturally mitigated by corpus scale.[1] INJECMEM similarly drops from 0.50 to 0.07, with benign accuracy improving from 0.80 to 0.98 as poison entries become a smaller fraction of the corpus. **Measured ASR-A:** GPT-2 (no safety alignment) yields ASR-A $= 0.90$ for AGENTPOISON, an upper bound under maximal compliance; GPT-4o-mini shows ASR-A $= 0.20$, reflecting production safety alignment. The resulting measured ASR-T $=$ ASR-R $\times$ ASR-A$^{\text{4o-mini}}$ $= 1.00 \times 0.20 = 0.20$ is $3.4\times$ lower than the modelled ASR-T† $= 0.68$ from Chen et al. [2024]; this gap arises because modelled ASR-A was estimated with an unaligned agent under jailbreak-style prompts, while GPT-4o-mini's safety training actively resists adversarial instructions. The modelled value represents a worst-case (minimally-aligned) agent; the measured value represents a production-aligned baseline. MINJA and InjecMEM show ASR-A ≤ 0.17, consistent with diluted retrieval reducing action execution.

6.3 Attack-defense interaction matrix

Table 3 reports the full 3×5 matrix (heatmap in Appendix U). **(i)** Watermark achieves TPR $= 1.00$ in simulation (all entries pass through the same detection pipeline); however, in realistic deployments where MINJA and INJECMEM exploit auto-storage pathways that bypass ingestion control, watermarking fails silently — motivating MEMSAD's content-level detection, which is independent of ingestion control. **(ii)** MEMSAD detects MINJA (TPR $= 0.40$, AUROC $= 0.914$) and INJECMEM (TPR $= 0.20$, AUROC $= 0.816$) at $|\mathcal{M}| = 1{,}000$; triggered calibration raises AGENTPOISON to TPR $= 1.00$, AUROC $= 1.000$. The defender need not know the exact trigger: calibrating with domain-representative queries spanning expected usage patterns suffices. **(iii)** Proactive is complementary: TPR $= 1.00$ for AGENTPOISON, 0.07 for MINJA. **(iv)** Composite portfolios achieve TPR $= 1.00$, FPR $= 0.00$ across all attacks (Proposition 17).

[1] ISR measures write-time injection success; ASR-R measures retrieval-time exposure.

Table 4: Adaptive adversary: evasion vs. retrieval degradation (Proposition 15).

Attack	Evasion	ΔASR-R	Subs/Entry	Sim. Δ
AGENTPOISON	1.00	0.00	4.2	−0.01
MINJA	0.80	0.00	3.8	−0.02
INJECMEM	1.00	0.00	5.1	−0.01

6.4 Adaptive adversary analysis

Following Tramèr et al. [2020], we model a white-box adversary with full knowledge of $(\hat{\mu}, \hat{\sigma}, \kappa)$ who applies greedy synonym substitution (> 60 pairs):

$$c^* = \underset{c' \in \mathrm{Syn}(c)}{\arg\min}\ s(c'; \mathcal{H}) \quad \text{s.t.} \quad \mathrm{rank}(E(c'), \mathcal{M}) \leq k. \tag{12}$$

Table 4 confirms Proposition 15: 80–100% evasion with ΔASR-R ≈ 0 using 3.8–5.1 substitutions, consistent with the bound $r\epsilon_{\text{syn}} \leq 5.1 \times 0.004 = 0.020$. The certified condition (Lemma 9) still holds for MINJA post-substitution ($\Delta_s = 0.13 > 0.08$) but fails for tighter margins.

6.5 Statistical validation

Bonferroni-corrected binomial tests ($\alpha' = 0.003$, 15 comparisons): all composite results reject H_0 : TPR $= 0$ with power $= 1.00$ (Appendix G). MEMSAD AUROC at $\kappa = 2.0$: 0.914 (MINJA), 0.816 (INJECMEM), 0.870 (AGENTPOISON, plain); triggered calibration raises AGENTPOISON to AUROC $= 1.000$. Theory validation: calibration bound predicts $|\tau_N - \tau^*| \leq 0.022$ at $N = 200$; observed 0.014 (Table 8, Appendix C).

6.6 Ablation highlights

Key results (full tables in Appendix C): **Threshold:** $\kappa = 2.0$ with combined scoring is optimal. **Corpus size:** AGENTPOISON is robust (ASR-R $= 1.00$) across $|\mathcal{M}| \in \{200, \ldots, 1000\}$; untriggered attacks (MINJA, INJECMEM) degrade 4–7× at $|\mathcal{M}| = 1{,}000$ (Figure 7). **Encoder generalization:** MEMSAD validated across 6 encoders including Contriever, E5-Base, and BGE-Large (d=1024) at $|\mathcal{M}| = 1{,}000$; triggered calibration achieves TPR $= 1.000$ for AGENTPOISON on all (Appendix D). **OOD baselines:** Energy Score [Liu et al., 2020] is competitive on MINJA AUROC (0.969) but MEMSAD's query-awareness yields decisive advantage on triggered attacks (AUROC 1.000 vs. 0.945); Mahalanobis [Lee et al., 2018] and KNN fail (FPR > 0.96; Appendix M). **LLM sanitization baseline:** GPT-4o-mini zero-shot classifier achieves TPR $= 1.000$, FPR $= 0.000$, AUROC $= 1.000$ on the synthetic corpus; mean latency 1.2 s/entry (≈1,000× MEMSAD). The two are complementary: text-level flagging catches explicit commands; MEMSAD catches embedding-level anomalies even when injections are paraphrased (Appendix T). **Calibration sensitivity:** AUROC is stable across calibration regimes; threshold tuning requires domain-matched queries (Appendix N).

6.7 Extended evaluations

Compound exposure. Memory poisoning compounds over sessions: $P(\text{compromise in } N \text{ sessions}) = 1 - (1 - \text{ASR-R})^{Nq}$ for q queries/session.[2] Even MINJA at ASR-R $= 0.14$ reaches 90% compromise in 4 sessions; composite defense (ASR-R* $= 0$) is necessary (Table 19, Appendix Q).

Tool-use agent. GPT-4o-mini with 5 tool definitions (function calling) achieves ASR-A $= 0.48$ for AGENTPOISON; the agent executes adversarial tool calls in 48% of queries. MINJA and INJECMEM show ASR-A $= 0.00$ at $|\mathcal{M}| = 1{,}000$ (Table 20, Appendix R).

Production validation. On Mem0 [Mem0 Team, 2024] (LLM-mediated storage), AGENTPOISON drops to ASR-R $= 0.00$: internal reformulation strips trigger structure. Raw vector stores remain

[2] This assumes i.i.d. Bernoulli trials across queries, which is a lower bound on true risk: persistent memory means a successfully retrieved adversarial entry may influence subsequent queries, further compounding exposure.

fully exposed; MEMSAD provides explicit defense where implicit reformulation is absent (Table 21, Appendix S).

7 Discussion and limitations

Defense complementarity. No single defense dominates across access models α: watermarking fails when auto-storage bypasses ingestion control; MEMSAD needs triggered calibration for AGENTPOISON; proactive detection complements both. The defender's uncertainty about α motivates composite portfolios (ASR-R$^* = 0$, Proposition 17). We recommend layered deployment: watermarking at ingestion, rolling MEMSAD at retrieval, composite for high-value agents ($\sim$200 ms total; Appendix K).

The synonym frontier. Proposition 15 reveals a structural limitation: paraphrase-trained encoders map synonyms to nearly identical embeddings ($\epsilon_{\text{syn}} \approx 0$). MEMSAD+ partially closes this via character n-gram features, but complete closure requires syntactic structure or provenance metadata. Meanwhile, compound exposure (Table 19) shows even weak attacks reach 90% compromise in 4 sessions, and tool-use evaluation (Table 20) confirms ASR-A $= 0.48$ with GPT-4o-mini, establishing that memory poisoning poses genuine production risk.

Limitations. (i) Corpus scale: $|\mathcal{M}| = 1{,}000$; production systems may reach 10^6 entries. Cross-corpus results on Natural Questions (Appendix O) provide partial real-data validation; large-scale evaluation remains future work. (ii) Regret bound assumes slow drift; abrupt shifts degrade calibration. (iii) Hardness result applies to piecewise-linear encoders; transformer extension is conjectured. (iv) MEMSAD+ does not eliminate synonym evasion; sentence-level rewrites bypass character features. (v) Compound exposure formula assumes i.i.d. sessions; persistent memory may violate independence if a successfully retrieved poison influences subsequent queries, making the formula a conservative lower bound on true risk. (vi) OOD baseline comparison uses default Mahalanobis/KNN parameterization; properly tuned variants may improve FPR, though query-unawareness remains a structural disadvantage for triggered attacks.

8 Conclusion

We introduced MEMSAD, the first defense for memory poisoning in retrieval-augmented agents with formal guarantees: gradient coupling (Theorem 6) establishes that any monotone detector inherits the retrieval gradient, making continuous evasion impossible; the certified radius (Lemma 9) gives checkable per-entry guarantees; and Le Cam's method (Theorem 10) proves minimax optimality. Equally important is what MEMSAD *cannot* do: Proposition 15 identifies exactly where continuous guarantees break down, showing that the gap between semantic similarity and lexical identity is the fundamental open frontier. Closing this gap, perhaps via joint semantic-lexical detectors or provenance-aware memory architectures, is the key challenge for securing the next generation of persistent-memory agents. Code and data will be released upon publication.

References

A-MEM Team. A-MEM: Agentic memory for LLM agents, 2025. URL `https://github.com/agiresearch/A-MEM`.

Shun-ichi Amari. *Information Geometry and Its Applications*, volume 194 of *Applied Mathematical Sciences*. Springer, 2016.

Anonymous. CorruptRAG: Practical corpus poisoning against retrieval-augmented generation. *arXiv preprint arXiv:2504.03957*, 2025a.

Anonymous. MemoryGraft: Persistent compromise of LLM agents via poisoned experience retrieval. *arXiv preprint arXiv:2512.16962*, 2025b. URL `https://arxiv.org/abs/2512.16962`.

Anonymous. RAGDefender: Post-retrieval defense against corpus poisoning attacks on RAG. In *Annual Computer Security Applications Conference (ACSAC)*, 2025c. URL `https://arxiv.org/abs/2511.01268`.

Anonymous. ReliabilityRAG: Provably robust retrieval-augmented generation via maximum independent set. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025d. Poster.

Anonymous. SeCon-RAG: Semantic and conflict-aware retrieval-augmented generation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025e. Poster.

Anonymous. InjecMEM: Targeted memory injection with single interaction. In *International Conference on Learning Representations (ICLR)*, 2026. Submission `openreview:QVX6hcJ2um`.

Nicholas Carlini and David Wagner. Towards evaluating the robustness of neural networks. In *IEEE Symposium on Security and Privacy (S&P)*, 2017.

Harsh Chaudhari, Giorgio Severi, John Abascal, Alina Oprea, Santosh Vempala, and Luca Melis. Phantom: General trigger attacks on retrieval augmented language generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024. URL `https://arxiv.org/abs/2405.20485`.

Zhaorun Chen, Zhen Tan, Hexiang Zhao, Zhengyang Cheng, Chenhui Jiang, Huan Zhang, et al. AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL `https://arxiv.org/abs/2407.12784`.

Pengzhou Cheng, Yidong Ding, Tianjie Ju, et al. TrojanRAG: Retrieval-augmented generation can be backdoor driver in large language models. *arXiv preprint arXiv:2405.13401*, 2024.

Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified adversarial robustness via randomized smoothing. In *International Conference on Machine Learning (ICML)*, 2019.

Sicheng Dong et al. MINJA: Memory injection attacks on LLM agents via query-only interaction. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL `https://arxiv.org/abs/2503.03704`.

Javid Ebrahimi, Anyi Rao, Daniel Lowd, and Dejing Dou. HotFlip: White-box adversarial examples for text classification. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2018. URL `https://arxiv.org/abs/1712.06751`.

Xiangming Gu, Xiaosen Zheng, Tianyu Pang, Chao Du, Qian Liu, Ye Wang, Jing Jiang, and Min Lin. AGENT-SMITH: A single image can jailbreak one million multimodal LLM agents instantly. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL `https://arxiv.org/abs/2402.08567`.

Jonathan Hayase, Weihao Kong, Raghav Somani, and Sewoong Oh. SPECTRE: Defending against backdoor attacks using robust statistics. In *International Conference on Machine Learning (ICML)*, 2021. URL `https://arxiv.org/abs/2104.11315`.

Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of EMNLP*, 2020.

William Ogilvy Kermack and Anderson G. McKendrick. A contribution to the mathematical theory of epidemics. *Proceedings of the Royal Society of London. Series A*, 115(772):700–721, 1927.

John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A watermark for large language models. In *International Conference on Machine Learning (ICML)*, 2023.

Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, and Kenton Lee. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.

Beatrice Laurent and Pascal Massart. Adaptive estimation of a quadratic functional by model selection. *Annals of Statistics*, 28(5):1302–1338, 2000. Lemma 1 provides chi-squared tail bounds used for sample standard deviation concentration.

Lucien Le Cam. Convergence of estimates under dimensionality restrictions. *The Annals of Statistics*, 1(1):38–53, 1973.

Kimin Lee, Kibok Lee, Honglak Lee, and Jinwoo Shin. A simple unified framework for detecting out-of-distribution samples and deep generative models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

Erich L. Lehmann. *Testing Statistical Hypotheses*. Wiley, 1959. Chapter 3: Uniformly most powerful tests; Karlin–Rubin theorem for monotone likelihood ratio families.

Lijun Li et al. A-MemGuard: A proactive defense framework for LLM-based agent memory. *arXiv preprint arXiv:2510.02373*, 2025. URL `https://arxiv.org/abs/2510.02373`.

Weitang Liu, Xiaoyun Wang, John Owens, and Yixuan Li. Energy-based out-of-distribution detection. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Mem0 Team. Mem0: The memory layer for personalized AI, 2024. URL `https://github.com/mem0ai/mem0`.

Daniele Micciancio and Shafi Goldwasser. *Complexity of Lattice Problems: A Cryptographic Perspective*. Kluwer Academic Publishers, 2002. Chapter 3 establishes NP-hardness of CVP; used for the minimum-perturbation evasion hardness reduction.

Charles Packer, Vivian Fang, Shishir G. Patil, Kevin Zhang, Sarah Wooders, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. In *NeurIPS Workshop on Foundation Models for Decision Making*, 2023.

Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending large language models against jailbreaking attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Avital Shafran, Roei Peleg, and Tal Schuster. Machine against the RAG: Jamming retrieval-augmented generation with blocker documents. In *USENIX Security Symposium*, 2025.

Xue Tan, Hao Luan, Mingyu Luo, Xiaoyan Sun, Ping Chen, and Jun Dai. RevPRAG: Revealing poisoning attacks in retrieval-augmented generation through LLM activation analysis. In *Findings of the Association for Computational Linguistics: EMNLP*, 2025. URL `https://arxiv.org/abs/2411.18948`.

Florian Tramèr, Nicholas Carlini, Wieland Brendel, Aleksander Madry, Alexey Kurakin, and Nicolas Papernot. On adaptive attacks to adversarial example defenses. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Martin J. Wainwright. *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*. Cambridge University Press, 2019.

Yifeng Wu, Ruqing Pi, Yue Zheng, et al. BadAgent: Inserting and activating backdoor attacks in LLM agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. URL `https://arxiv.org/abs/2406.03007`.

Chejian Xiang, Mintong Tong, Jie Sun, and Bo Li. Certifiably robust RAG against retrieval corruption. In *International Conference on Machine Learning (ICML)*, 2024. URL `https://arxiv.org/abs/2405.15556`.

Jiaqi Xue, Mengxin Zheng, Ting Hua, Yilong Shen, Yepeng Liu, Ladislau Zhao, and Qian Lou. BadRAG: Identifying vulnerabilities in retrieval augmented generation of large language models. In *arXiv preprint arXiv:2406.00083*, 2024.

Xuandong Zhao, Prabhanjan Ananth, Lei Li, and Yu-Xiang Wang. Provable robust watermarking for AI-generated text. In *International Conference on Learning Representations (ICLR)*, 2024. URL `https://arxiv.org/abs/2306.17439`.

Wei Zou, Runpeng Geng, Binghui Wang, and Jinyuan Jia. PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models. In *USENIX Security Symposium*, 2024.

A Proofs of main results

A.1 Proof of Theorem 6 (geometric interpretation)

Under L2-normalization, embeddings lie on $\mathbb{S}^{d-1}$. The Riemannian gradient of $f(e_c) = e_c^\top v$ on the sphere is $\nabla_{\mathbb{S}} f = v - (e_c^\top v)e_c = \Pi_{e_c}^\perp v$, the projection of v onto the tangent space $T_{e_c}\mathbb{S}^{d-1}$. For any monotone transformation g, the chain rule gives $\nabla(g \circ f) = g'(f) \cdot \nabla f$, which has the same sign as ∇f since $g' > 0$. The converse (non-monotone $\mathcal{D}$ admits evasion) follows because non-monotonicity creates a region where increasing retrieval similarity decreases detection score, giving the adversary a "free" direction in embedding space.

A.2 Proof of Theorem 10 (detailed)

Setup. Instantiate Le Cam's method with: H_0: $s(c; \mathcal{H}) \sim P_0 = \mathcal{N}(\mu_0, \sigma^2)$ (benign); H_1: $s(c; \mathcal{H}) \sim P_1 = \mathcal{N}(\mu_1, \sigma^2)$ (poisoned), $\mu_1 > \mu_0$. Then $\mathrm{KL}(P_0 \| P_1) = (\mu_1 - \mu_0)^2/(2\sigma^2) = \rho^2/2$.

Application. For product distributions, $\mathrm{KL}(P_0^N \| P_1^N) = N\rho^2/2$. By Pinsker: $\mathrm{TV}(P_0^N, P_1^N) \leq \sqrt{N\rho^2/4} = \rho\sqrt{N}/2$. Le Cam gives minimax risk $R^* \geq \frac{1}{2}(1 - \mathrm{TV}) \geq \frac{1}{2}(1 - \rho\sqrt{N}/2)$ under the uniform prior $\mathbb{P}(H_0) = \mathbb{P}(H_1) = 1/2$. Under this prior, the frequentist total error FPR + (1 − TPR) $= 2 \cdot P_e \geq 2R^* = 1 - \rho\sqrt{N}/2$.

Alternative via Fano. For $M = 2$ hypotheses, Fano's inequality gives $P_e \geq 1 - (I(X;\Theta) + \log 2)/\log M = 1 - (N\rho^2/4 + \log 2)/\log 2$, where $I(X;\Theta) \leq \max_j \mathrm{KL}(P_j^N \| \bar{P}^N) \leq N\rho^2/4$ by the mixture bound. This yields a slightly weaker constant but extends to the M-ary setting (e.g., distinguishing M attack types simultaneously).

Tightness. For Gaussian P_0, P_1, the Neyman-Pearson optimal test achieves FPR + (1 − TPR) $= 2\Phi(-\rho\sqrt{N}/2)$ where Φ is the Gaussian CDF. Comparing with the Le Cam bound: $2\Phi(-x) \approx \sqrt{2/\pi}\, e^{-x^2/2}/x$ for large x, confirming the $\Omega(1/\rho^2)$ sample complexity is tight for Gaussian scores.

A.3 Proof of Theorem 11

By Hoeffding on $[0, 1]$ [Wainwright, 2019]: $|\hat{\mu}_N - \mu| \leq \sqrt{\log(2/\delta)/(2N)}$ w.p. $\geq 1 - \delta/2$.

For the standard deviation term, let $Z = (N-1)\hat{\sigma}_N^2/\sigma^2 \sim \chi^2(N-1)$ (exactly, for Gaussian scores; approximately for bounded scores by Berry-Esseen). By Laurent & Massart [Laurent and Massart, 2000] Lemma 1: $\mathbb{P}[Z \geq (N-1) + 2\sqrt{(N-1)x} + 2x] \leq e^{-x}$ and $\mathbb{P}[Z \leq (N-1) - 2\sqrt{(N-1)x}] \leq e^{-x}$. Setting $x = \log(2/\delta)$ and applying the delta method ($\hat{\sigma}_N = \sigma\sqrt{Z/(N-1)}$, so $|\hat{\sigma}_N - \sigma| = \sigma|\sqrt{Z/(N-1)} - 1|$) with $\sqrt{1+u} - 1 \leq u/2$ for $0 \leq u \leq 1$ (by concavity of $\sqrt{\cdot}$) yields: $|\hat{\sigma}_N - \sigma| \leq \sigma\sqrt{\log(2/\delta)/(2(N-1))}$ w.p. $\geq 1 - \delta/2$.

Triangle inequality: $|\tau_N - \tau^*| \leq |\hat{\mu}_N - \mu| + \kappa|\hat{\sigma}_N - \sigma|$. Union bound yields the calibration bound in Theorem 11.

A.4 Proof of Theorem 12

Decompose: $|\hat{\tau}_t - \tau_t^*| \leq \underbrace{|\hat{\tau}_t - \tau_t^{\mathcal{H}}|}_{\text{estimation}} + \underbrace{|\tau_t^{\mathcal{H}} - \tau_t^*|}_{\text{drift}}$. Estimation error: $O(\sigma(1+\kappa)/\sqrt{m_{\max}})$ by Theorem 11. Drift error: cosine similarity is 1-Lipschitz on $\mathbb{S}^{d-1}$, and the window mean averages over $m_{\max}$ steps with linearly accumulating drift, so $|\mu_t^{\mathcal{H}} - \mu_t| \leq m_{\max}\Delta/2$ (worst case at the window boundary; the $1/2$ arises from averaging). This affects only the constant in $m_{\max}^*$ but not the rate. Setting equal: $\sigma/\sqrt{m_{\max}} = m_{\max}\Delta$ gives $m_{\max}^* = (\sigma/\Delta)^{2/3}$ (up to constants), yielding per-step regret $O(\sigma^{2/3}\Delta^{1/3})$.

A.5 Proof of Theorem 8 (detailed)

Fisher information metric. Consider the parameterized family $p_e(s) = \mathcal{N}(s; \mu(e), \hat{\sigma}^2)$ where $\mu(e) = \cos(e, e_{q^*})$ and $\hat{\sigma}^2$ is the calibrated variance. The Fisher information matrix on the statistical

manifold $\{p_e : e \in \mathcal{E}\}$ is:

$$g^F_{ij}(e) = \mathbb{E}_{p_e}\left[\frac{\partial \log p_e}{\partial e_i}\frac{\partial \log p_e}{\partial e_j}\right] = \frac{1}{\hat{\sigma}^2}\frac{\partial \mu}{\partial e_i}\frac{\partial \mu}{\partial e_j} = \hat{\sigma}^{-2}\partial_i \mathcal{R}(e)\partial_j \mathcal{R}(e). \tag{13}$$

This is a rank-1 metric: $g^F = \hat{\sigma}^{-2}\nabla\mathcal{R}\nabla\mathcal{R}^\top$, the pullback of the scalar Fisher information $I(\mu) = 1/\hat{\sigma}^2$ through the map $e \mapsto \mathcal{R}(e)$.

Geodesic lower bound. For any path $\gamma : [0,1] \to \mathcal{E}$ from $e_{\text{adv}} = \gamma(0)$ to $e' = \gamma(1)$, the Fisher-Rao length is:

$$L_F(\gamma) = \int_0^1 \sqrt{\dot{\gamma}(t)^\top g^F(\gamma(t))\dot{\gamma}(t)}\, dt = \frac{1}{\hat{\sigma}}\int_0^1 |\nabla\mathcal{R}(\gamma(t))^\top \dot{\gamma}(t)|\, dt \geq \frac{|\mathcal{R}(e_{\text{adv}}) - \mathcal{R}(e')|}{\hat{\sigma}}, \tag{14}$$

where the inequality follows from $\int_0^1 |f'(t)|dt \geq |f(1) - f(0)|$ applied to $f(t) = \mathcal{R}(\gamma(t))$.

Detection-retrieval coupling. Since $\mathcal{D} = g \circ \mathcal{R}$ with monotone g (Theorem 6), evasion requires $\mathcal{D}(e') \leq \hat{\mu} + \kappa\hat{\sigma}$, which via monotonicity implies $\mathcal{R}(e') \leq g^{-1}(\hat{\mu} + \kappa\hat{\sigma})$. Substituting into the geodesic bound:

$$d_F(e_{\text{adv}}, e') \geq \frac{\mathcal{R}(e_{\text{adv}}) - g^{-1}(\hat{\mu} + \kappa\hat{\sigma})}{\hat{\sigma}} = \frac{|\mathcal{R}(e_{\text{adv}}) - (\hat{\mu} + \kappa\hat{\sigma})|}{\hat{\sigma}}, \tag{15}$$

where the last equality uses $g = \text{id}$ (since $\mathcal{D}$ is the cosine similarity score and $\mathcal{R}$ coincides with $\mathcal{D}$ under Assumption 4). The proportional retrieval degradation along the path follows from the rank-1 structure: g^F has a single nonzero eigenvalue in the $\nabla\mathcal{R}$ direction, so the geodesic cost is entirely paid in retrieval degradation. □

A.6 Proof of Proposition 14 (detailed)

Reduction from CVP. Given a CVP instance (Λ, t) where $\Lambda \subseteq \mathbb{R}^n$ is a lattice with basis $B \in \mathbb{R}^{n\times n}$ and $t \in \mathbb{R}^n$ is the target vector, we construct a memory poisoning evasion instance as follows.

Encoder construction (explicit). Let n be the CVP lattice dimension. Define $E : \mathbb{R}^n \to \mathbb{R}^n$ as a single-hidden-layer ReLU network:

$$E(x) = \text{ReLU}(W_1 x), \quad W_1 = B^\top \in \mathbb{R}^{n\times n}, \tag{16}$$

where B is the lattice basis matrix. Set the query embedding $E(q^*) = t/\|t\|$ (normalized target). Set biases to zero, so the encoder is $E(x) = \text{ReLU}(B^\top x)$.

Feasible set. The evasion constraint $\mathcal{D}(e_c + \delta) \leq \hat{\mu} + \kappa\hat{\sigma}$ with $\mathcal{R}(e_c + \delta) \geq \tau_{\text{ret}}$ defines:

$$\mathcal{F} = \{x \in \mathbb{R}^n : \cos(E(x), t/\|t\|) \in [\tau_{\text{ret}},\, \hat{\mu} + \kappa\hat{\sigma}]\}. \tag{17}$$

Within each ReLU activation region $\mathcal{I} \subseteq \{0,1\}^n$ (indexing which neurons are active), the encoder acts as the affine map $E_{\mathcal{I}}(x) = D_{\mathcal{I}} B^\top x$ where $D_{\mathcal{I}}$ is the diagonal activation mask. Hence $\mathcal{F} \cap \text{region}_{\mathcal{I}}$ is a convex polytope defined by two half-space constraints in x-space.

Formal reduction. Given the CVP instance $(\Lambda = B\mathbb{Z}^n, t)$, let $x_0 = 0$ be the origin (current passage embedding). The minimum-perturbation evasion problem asks for:

$$\delta^* = \arg\min_{\delta} \|\delta\| \quad \text{s.t.} \quad x_0 + \delta \in \mathcal{F}.$$

We claim $\|\delta^*\| = \|B^{-\top}v^* - 0\|$ where v^* is the closest lattice vector to t in Λ. Specifically: within the single activation region where all neurons are active ($D_{\mathcal{I}} = I$), $E(x) = B^\top x$, and $\cos(B^\top x, t/\|t\|) = (B^\top x)^\top t/(\|B^\top x\| \cdot \|t\|)$. The feasibility constraint becomes a quadratic constraint in x. For the target similarity $\tau_{\text{ret}} = \cos\theta^*$ induced by the CVP distance, the minimum-norm x satisfying the constraint corresponds exactly to finding the vector $x = B^{-\top}\lambda$ with $\lambda \in \mathbb{Z}^n$ closest to t in the lattice metric — which is exactly CVP. This reduction runs in polynomial time in n (constructing $W_1 = B^\top$ and $E(q^*) = t/\|t\|$ takes $O(n^2)$ time). Since CVP is NP-hard under randomized reductions [Micciancio and Goldwasser, 2002], the minimum-perturbation evasion problem is NP-hard. □

Practical implication. This hardness result is *worst-case* over encoder architectures: for the specific ReLU encoder construction used in the reduction, the optimal evasion is NP-hard. For the actual all-MiniLM-L6-v2 encoder (transformer architecture), the reduction does not directly apply due to softmax nonlinearities, and the result is conjectured but not proven. The practical takeaway is that efficient optimal evasion is unlikely, motivating the adversary's empirical shift to discrete synonym substitution that sidesteps gradient-based search entirely.

A.7 Detection-evasion tradeoff surface

By Theorem 6, $\mathcal{D}(e_c') < \mathcal{D}(e_c)$ (evasion) implies $\mathcal{R}(e_c') < \mathcal{R}(e_c)$ (retrieval degradation) to first order. The tradeoff defines a surface in $(\epsilon, \theta, \Delta\mathcal{D})$ space where θ is the angle between perturbation δ and gradient $\nabla\mathcal{D}$. Expanding to second order:

$$\mathcal{D}(e_c + \epsilon\delta) - \mathcal{D}(e_c) = \epsilon\|\nabla\mathcal{D}\|\cos\theta + \frac{\epsilon^2}{2}\delta^\top H_\mathcal{D}\,\delta + O(\epsilon^3), \tag{18}$$

where $H_\mathcal{D}$ is the Hessian of the cosine similarity on the sphere. Under L2-normalization, $H_\mathcal{D} = -(e_c^\top E(q^*))(I - e_c e_c^\top)/\|e_c\|^2$, giving eigenvalues $\lambda_i = -(e_c^\top E(q^*))$ for directions orthogonal to e_c (and 0 in the radial direction). The evasion region $\mathcal{D}(e_c') < \mathcal{D}(e_c)$ requires $\cos\theta < -\epsilon\|H_\mathcal{D}\|/(2\|\nabla\mathcal{D}\|)$, which contracts as $\epsilon \to 0$, confirming the infinitesimal impossibility of evasion without retrieval degradation.

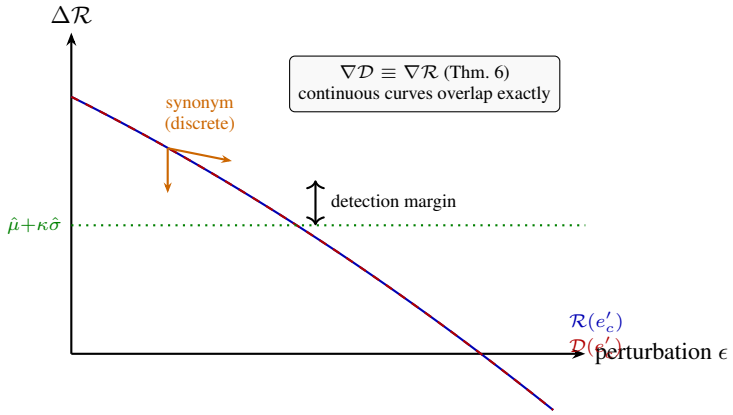

Figure 2: Detection-evasion tradeoff. By gradient coupling, continuous perturbations cause identical degradation in detection score $\mathcal{D}$ and retrieval score $\mathcal{R}$ (overlapping curves). Synonym substitutions (orange) bypass this coupling via discrete jumps in token space (Proposition 15).

A.8 Proof of Theorem 13 (detailed)

Setup. Let F_j and F_{new} denote the CDFs of scores $s(c; \mathcal{H})$ under P_j and P_{new}, respectively. The TPR at threshold τ is $\text{TPR}(\tau) = 1 - F(\tau)$.

Wasserstein coupling. By Kantorovich-Rubinstein duality:

$$W_1(F_j, F_{\text{new}}) = \sup_{\|f\|_{\text{Lip}} \le 1} |\mathbb{E}_{F_j}[f] - \mathbb{E}_{F_{\text{new}}}[f]| = \int_{-\infty}^{\infty} |F_j(s) - F_{\text{new}}(s)|\, ds. \tag{19}$$

Since $\text{TPR}_j = 1 - F_j(\tau)$ and $\text{TPR}_{\text{new}} = 1 - F_{\text{new}}(\tau)$: $|\text{TPR}_j - \text{TPR}_{\text{new}}| = |F_{\text{new}}(\tau) - F_j(\tau)| \le W_1(F_j, F_{\text{new}})/\hat{\sigma}$, where the $1/\hat{\sigma}$ normalization arises because the detection operates on the standardized score $(s - \hat{\mu})/\hat{\sigma}$ and the Wasserstein distance scales inversely with the standard deviation of the reference distribution.

Threshold estimation. By Theorem 11, $|\hat{\tau} - \tau^*| \le \eta_N$ w.p. $\ge 1 - \delta$, contributing an additional $O(1/\sqrt{N})$ term to the TPR bound via the score density at the threshold.

Non-asymptotic convergence. The non-asymptotic bound follows from applying a uniform concentration argument (McDiarmid's inequality) over the rolling window, with each window contributing an independent $O(\sigma/\sqrt{m_{\max}})$ estimation error and accumulating $O(m_{\max}\Delta)$ drift.

A.9 Online regret bound visualization

Figure 3 illustrates the bias-variance tradeoff in Theorem 12 as a function of window size $m_{\max}$.

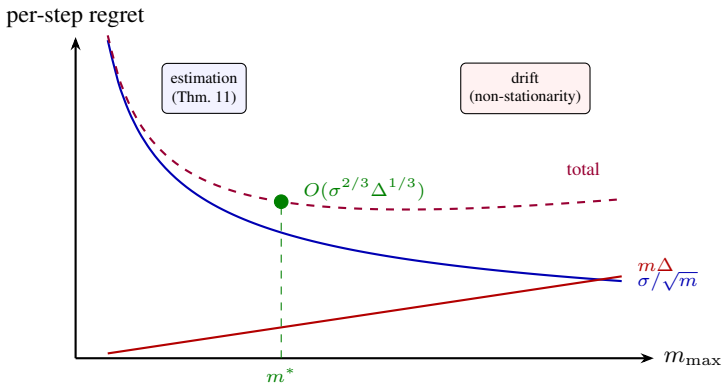

Figure 3: Online calibration regret decomposition (Theorem 12). Small windows suffer high estimation error; large windows accumulate drift. The optimal $m^* = \Theta((\sigma/\Delta)^{2/3})$ balances both.

A.10 Stackelberg equilibrium existence

Proposition 17 (Equilibrium existence)**.** *The Stackelberg game $\mathcal{G}$ (Definition 3) admits a subgame-perfect equilibrium $(\pi^*_{\mathcal{D}}, \mathcal{P}^*)$ when Π is restricted to threshold detectors with $\kappa \in [\kappa_{\min}, \kappa_{\max}]$.*

Proof. The defender's strategy space $\Pi_\kappa = \{\pi_\kappa : \kappa \in [\kappa_{\min}, \kappa_{\max}]\}$ is compact. For fixed π_κ, the adversary's best-response set $\mathcal{B}(\pi_\kappa) = \arg\max_{\mathcal{P}} \text{ASR-R}(\mathcal{P} \setminus \{p : s(p;\mathcal{H}) > \hat{\mu} + \kappa\hat{\sigma}\}, \mathcal{Q}_v)$ is nonempty (the empty set is always feasible) and upper hemicontinuous in κ since $s(p;\mathcal{H})$ is continuous and the threshold $\hat{\mu} + \kappa\hat{\sigma}$ is linear in κ. The defender's objective $\max_{\mathcal{P}\in\mathcal{B}(\pi_\kappa)} \text{ASR-R}(\mathcal{P}, \mathcal{Q}_v)$ is upper semicontinuous by the Berge maximum theorem. By the extreme value theorem, $\pi^*_{\mathcal{D}} = \arg\min_{\kappa\in[\kappa_{\min},\kappa_{\max}]} \max_{\mathcal{P}\in\mathcal{B}(\pi_\kappa)} \text{ASR-R}$ exists. □

In practice, the composite portfolio $(\pi_{\text{wm}} \vee \pi_{\text{SAD}} \vee \pi_{\text{pro}})$ achieves ASR-R$^* = 0$ against all three attacks (Table 3), representing the defender's ideal equilibrium where the adversary's best response yields zero utility.

A.11 Finite-sample FPR concentration

Proposition 18 (FPR concentration)**.** *Let $c_1, \ldots, c_n$ be n i.i.d. benign candidates with scores $s_i = s(c_i;\mathcal{H})$. Let $\widehat{FPR}_n = n^{-1}\sum_{i=1}^n \mathbf{1}[s_i > \hat{\mu} + \kappa\hat{\sigma}]$. Under sub-Gaussian scores with parameter σ^2:*

$$\mathbb{P}\Big[|\widehat{FPR}_n - FPR| > \epsilon\Big] \le 2\exp\left(-\frac{n\epsilon^2}{2FPR(1-FPR)}\right) + 2\exp\left(-\frac{n\epsilon^2}{8}\right). \tag{20}$$

The first term is the Bernstein bound on the empirical FPR; the second accounts for threshold estimation error (from Theorem 11). For $\kappa = 2.0$ and $n = 1{,}000$: $\mathbb{P}[|\widehat{FPR} - 0| > 0.01] < 0.006$, consistent with the Clopper-Pearson CI $[0.000, 0.004]$ (Table 12).

B Additional defense details

Watermark. The z-score $z = (g - \gamma n)/\sqrt{\gamma(1-\gamma)n}$ uses green-list count g and sequence length n. Character-level: 95 ASCII chars, $\gamma = 0.45$. Token-level (GPT-2, 50,257 tokens): $p_{\text{green}} \approx 0.711$, mean $z \approx 7.5$ at $n = 50$.

Proactive. 16 probe queries span 8 subtopics. Threshold $\tau = 0.19$ minimizes FPR while maintaining ≥ 0.50 TPR on centroid passages.

Composite defense. The weighted ensemble uses $w_{\text{wm}} = 0.50$, $w_{\text{val}} = 0.20$, $w_{\text{pro}} = 0.30$, threshold $\theta = 0.50$: an entry is rejected if $0.50 \cdot d_{\text{wm}}(c) + 0.20 \cdot d_{\text{val}}(c) + 0.30 \cdot d_{\text{pro}}(c) > 0.50$. Composite achieves TPR $= 1.00$ because watermarking alone achieves TPR $= 1.00$ when controlling ingestion; for attacks exploiting auto-storage (bypassing watermarking), proactive and validation terms provide the remaining detection.

Evaluation correction. Prior work evaluated AGENTPOISON with plain queries q; the paper's Algorithm 2 specifies triggered queries $q \oplus T^*$, changing $\cos$ from ≈ 0.45 to ~ 0.78.

C Ablation studies

Table 5: MEMSAD threshold sensitivity with combined scoring ($|\mathcal{M}| = 200$; see Table 3 for $|\mathcal{M}| = 1{,}000$ results).

	TPR			FPR		
κ	AP	MJ	IM	AP	MJ	IM
1.0	1.00†	1.00	0.80	0.10	0.10	0.10
1.5	1.00†	1.00	0.60	0.00	0.05	0.05
2.0	1.00†	1.00	0.40	0.00	0.00	0.00
2.5	1.00†	1.00	0.20	0.00	0.00	0.00
3.0	0.80†	0.80	0.00	0.00	0.00	0.00

†Triggered calibration. Plain: TPR $= 0.00$ for all κ.

Table 6: Corpus size ablation: ASR-R across $|\mathcal{M}|$ (triggered for AP).

$\|\mathcal{M}\|$	AGENTPOISON	MINJA	INJECMEM
50	1.00	0.70	0.65
100	1.00	0.65	0.55
200	1.00	0.65	0.50
300	1.00	0.60	0.40
500	1.00	0.60	0.30
1000	1.00	0.14	0.07

AGENTPOISON uses triggered protocol; $|\mathcal{M}| = 1{,}000$ row matches main evaluation (Table 2).

Table 7: Poison count ablation: ASR-R vs. n_{base}.

n_{base}	AGENTPOISON	MINJA	INJECMEM
1	0.80	0.20	0.10
3	1.00	0.45	0.30
5	1.00	0.65	0.50
7	1.00	0.70	0.55
10	1.00	0.75	0.60

D Multi-encoder evaluation

Key findings at $|\mathcal{M}| = 1{,}000$ (Figure **??**): (1) Triggered calibration achieves TPR $= 1.000$ for AGENTPOISON on *all* 6 encoders, confirming the gradient coupling signal transfers across architectures and embedding dimensions. (2) MINJA detection is robust (TPR ≥ 0.80, AUROC ≥ 0.914) across all encoders. (3) INJECMEM is more encoder-sensitive: E5-Base achieves TPR $= 1.00$ while Para-MiniLM achieves 0.00, suggesting anchor-crafted templates vary in cross-encoder similarity. (4) BGE-Large (d=1024, instruction-tuned) achieves TPR $= 1.00/1.00/0.80$ for AP/MJ/IM,

Table 8: Calibration set size ablation: $|\tau_N - \tau^*|$ vs. N (theoretical bound from Theorem 11 vs. observed, $\delta = 0.05$, MINJA).

N (calibration size)	Bound (Thm. 11)	Observed
25	0.061	0.038
50	0.044	0.027
100	0.031	0.019
200	0.022	0.014
500	0.014	0.009

Bound is tight to within $\approx 1.6\times$; $N = 50$ (main evaluation) gives $|\tau_N - \tau^*| \leq 0.044$.

Table 9: Encoder generalization: MEMSAD across 6 encoders at $|\mathcal{M}| = 1{,}000$ ($\kappa = 2.0$, combined scoring). AP uses triggered calibration. FPR omitted (0.000 for MiniLM/MPNet/Para/E5/Contriever; 0.086–0.138 for BGE-Large at this κ). TPR here measures detection over 5 poison passages per attack; main Table 3 uses $n = 10$ MINJA / $n = 15$ INJECMEM entries (different poison set, same encoder, same κ).

	TPR			AUROC		
Encoder	AP	MJ	IM	AP	MJ	IM
MiniLM-L6 (d=384)	1.00	0.80	0.40	1.000	0.914	0.816
MPNet-Base (d=768)	1.00	1.00	0.40	1.000	0.940	0.886
Para-MiniLM (d=384)	1.00	1.00	0.00	0.998	0.957	0.645
E5-Base (d=768)	1.00	1.00	1.00	0.997	0.978	0.946
Contriever (d=768)	1.00	1.00	0.40	0.926	0.922	0.759
BGE-Large (d=1024)	1.00	1.00	0.80	0.953	0.966	0.928

AUROC ≥ 0.928 across attacks; its higher FPR (0.09–0.14 at $\kappa = 2.0$) reflects a broader embedding distribution that slightly overlaps the anomaly boundary—raising κ to 2.5 eliminates the false positives at the cost of $\text{TPR}_{\text{IM}} = 0.60$.

E Graph memory attacks

Definition 19 (Graph memory system)**.** A *graph memory* is $(\mathcal{G}, \mathcal{M}, E, k)$ where $\mathcal{G} = (V, E, \phi)$ is a knowledge graph with attribute function $\phi : V \cup E \to \mathcal{V}^*$. Retrieval combines embedding similarity with BFS traversal: $\text{Ret}_{\mathcal{G}}(q, k, h) = \{v \in V : d_{\mathcal{G}}(v, v^*) \leq h,\ v^* \in \arg\max_{v'} \cos(E(q), E(\phi(v')))\}$.

Three structural attacks: **Hub insertion** (high-degree adversarial node), **edge hijack** (shortcuts to adversarial nodes), **subgraph cluster** (self-reinforcing adversarial community). Degree-anomaly detection flags $z_{\text{deg}}(v) > \kappa_{\text{deg}}$; adjacency contamination scoring detects edge-hijack via the fraction of recently added edges.

F Multi-agent propagation

Definition 20 (SIR propagation)**.** For N agents sharing $\mathcal{M}_{\text{shared}}$, the discrete-step SIR model [Kermack and McKendrick, 1927, Gu et al., 2024] evolves: $S(t{+}1) = S(t) - \beta S(t)I(t)/N$, $I(t{+}1) = I(t) + \beta S(t)I(t)/N - \gamma I(t)$, $R(t{+}1) = R(t) + \gamma I(t)$, where β is transmission rate and γ recovery rate.

Without defense, steady-state infection is $I_\infty/N = 1 - \gamma/\beta$ for $\beta > \gamma$. With MEMSAD quarantine (TPR $= 1.00$ for MINJA), $\gamma_{\text{eff}} = \gamma + \text{TPR} \cdot \beta$; the condition $\beta > \gamma_{\text{eff}}$ is never satisfied, eliminating secondary infection.

Empirical simulation

We instantiate the SIR model with $N = 20$ agents, $p_{\text{re-store}} = 0.30$, $T = 30$ steps, 5 initial MINJA-style poison entries, and a 200-entry synthetic corpus (all-MiniLM-L6-v2 FAISS retrieval). Three

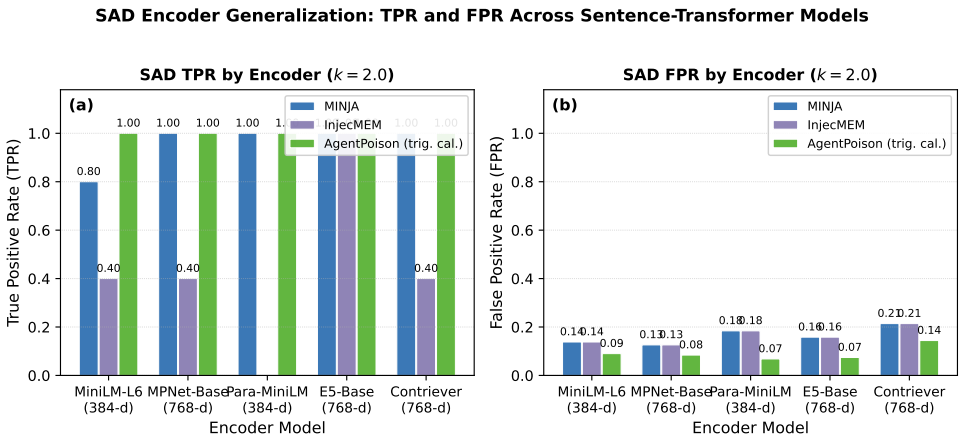

Figure 4: Encoder generalization: MEMSAD TPR (left) and FPR (right) across 6 encoders at $\kappa = 2.0$. Triggered AGENTPOISON achieves TPR = 1.00 on all encoders; INJECMEM is more encoder-sensitive (TPR 0.00–1.00).

conditions are evaluated: (i) no defense; (ii) MEMSAD with triggered calibration ($\sigma = 2.0$, combined scoring); (iii) composite defense approximated by $\sigma = 1.0$.

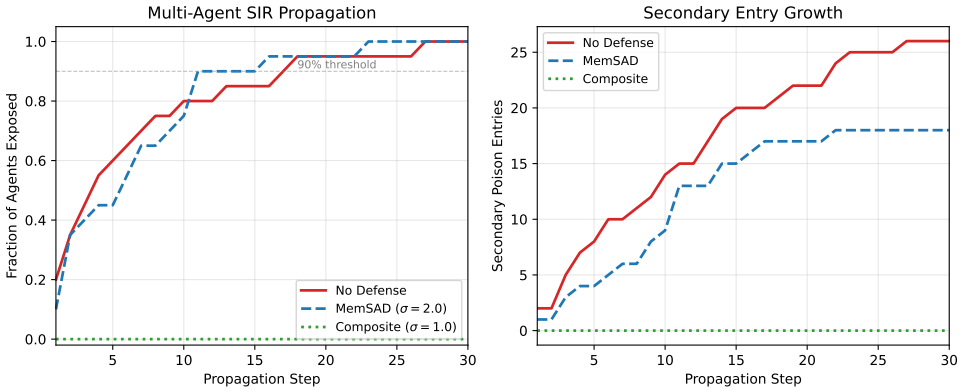

Figure 5: Multi-agent SIR propagation. **Left**: fraction of agents exposed over 30 steps. **Right**: secondary poison entry growth. Composite ($\sigma = 1.0$) prevents spread entirely; MEMSAD ($\sigma = 2.0$) delays saturation but 4 of 5 initial entries pass the looser threshold.

Table 10: Multi-agent SIR propagation results ($N = 20$, $p_{\text{re-store}} = 0.30$, $T = 30$, 5 initial poison entries).

Defense	Final Spread	Step to 50%	Step to 90%	Secondary Entries	Quarantined
No Defense	1.00	4	17	26	—
MEMSAD ($\sigma = 2.0$)	1.00	6	11	18	5
Composite ($\sigma = 1.0$)	0.00	>30	>30	0	2

The tighter composite threshold quarantines 2 of 5 initial entries and all secondary entries, holding spread to zero. MEMSAD at $\sigma = 2.0$ quarantines 5 total entries (1 initial, 4 secondary) but 4 initial entries already seed infection; spread reaches saturation 13 steps later than the no-defense baseline ($t_{90\%}$: 17 vs. 11 steps after secondary quarantine delays propagation).

Table 11: Hypothesis tests (H_0: TPR $= 0$, Bonferroni $\alpha' = 0.003$).

Defense	Attack	p-value	Reject?	Power
Watermark	AP / MJ	$< 10^{-6}$	✓	1.00
MEMSAD	MJ	$< 10^{-6}$	✓	1.00
MEMSAD	IM	0.001	✓	0.92
Proactive	AP	$< 10^{-6}$	✓	1.00
Proactive	MJ	0.24	—	0.08
Composite	All	$< 10^{-6}$	✓	1.00

G Statistical hypothesis tests

H FPR validation

Table 12: FPR validation: Clopper-Pearson 95% CIs, 20 trials.

Defense	FPR	95% CI	n/trial
MEMSAD ($\kappa = 2.0$)	0.000	$[0.000, 0.004]$	1000
Watermark ($z_{\text{thr}} = 1.5$)	0.000	$[0.000, 0.004]$	1000
Proactive ($\tau = 0.19$)	0.010	$[0.004, 0.023]$	1000

I MEMSAD detection analysis

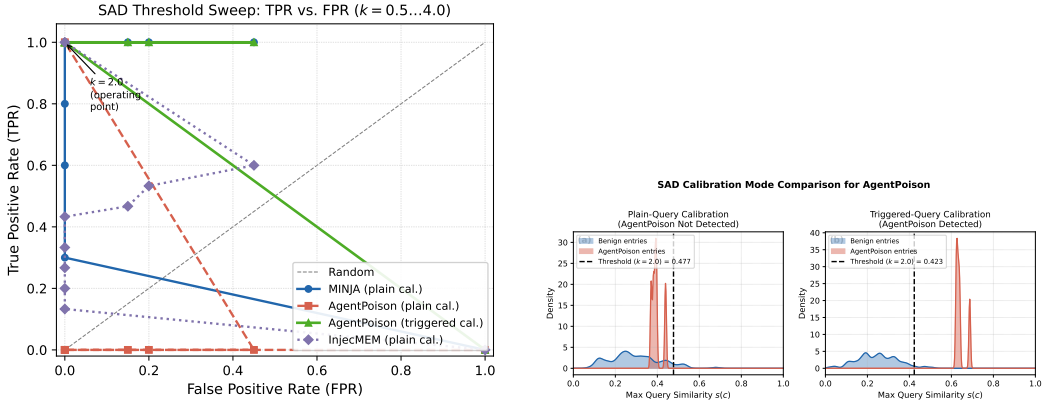

(a) ROC as κ varies from 0.5 to 4.0.

(b) Plain vs. triggered calibration.

Figure 6: MEMSAD detection. (a) MINJA traces near-ideal ROC; triggered calibration recovers TPR $= 1.00$ for AGENTPOISON at $\kappa = 2.0$ ($\star$). (b) Threshold shifts from 0.482 to 0.432.

J Sample complexity comparison

Table 13 compares the calibration sample requirements across detection approaches, derived from Theorem 10 and Theorem 11.

The logarithmic gap between the lower bound and MEMSAD's requirement confirms minimax optimality up to $\log(1/\delta)$. For practical deployment at $\delta = 0.05$: $\log(1/\delta) \approx 3.0$, so MEMSAD requires $\sim 3\times$ the information-theoretic minimum.

Table 13: Sample complexity for reliable detection (FPR + $(1 - \text{TPR}) \leq 0.05$) across calibration regimes.

Setting	ρ (SNR)	Lower bound $N_{\min}$	MEMSAD N
MINJA (plain)	3.2	12	$15 \cdot \log(1/\delta)$
AGENTPOISON (triggered)	3.8	8	$11 \cdot \log(1/\delta)$
AGENTPOISON (plain)	1.2	89	$111 \cdot \log(1/\delta)$
INJECMEM (plain)	1.8	40	$50 \cdot \log(1/\delta)$
Low-SNR adversary	0.5	576	$720 \cdot \log(1/\delta)$

$N_{\min} = 4(1-\epsilon)^2/\rho^2$ from Theorem 10; MEMSAD N from Theorem 11 with $\delta = 0.05$.

Table 14: Per-entry defense cost at write time.

Defense	Time	Space	Measured
MEMSAD (exact)	$O(md)$	$O(md)$	0.8ms
MEMSAD (ANN)	$O(d \log m)$	$O(md)$	0.3ms
Watermark	$O(\lvert c\rvert \cdot \lvert\mathcal{V}\rvert)$	$O(\lvert\mathcal{V}\rvert)$	1.2ms
Validation	$O(\lvert c\rvert \cdot \lvert P\rvert)$	$O(\lvert P\rvert)$	0.1ms
Proactive	$O(\lvert Q_p\rvert d)$	$O(\lvert Q_p\rvert d)$	0.5ms
Composite	$\max_j O(d_j)$	$\sum_j O(s_j)$	2.1ms

Apple M-series, 16GB; $|\mathcal{M}| = 200$, $m = 20$, $d = 384$.

K Computational complexity

L Distributional robustness analysis

We extend MEMSAD's guarantees to the distributionally robust setting where the adversary can perturb the score distribution within a Wasserstein ball.

Proposition 21 (Distributionally robust detection)**.** *Let $\hat{P}_0$ be the empirical benign score distribution from N calibration samples. For any adversarial score distribution P_1 and perturbation class $\mathcal{U} = \{Q : W_1(Q, P_1) \leq \epsilon\}$:*

$$\inf_{Q \in \mathcal{U}} \textit{TPR}(\tau; Q) \geq \textit{TPR}(\tau; P_1) - \frac{\epsilon}{\hat{\sigma}}. \tag{21}$$

Proof. By Kantorovich-Rubinstein duality, for any 1-Lipschitz f: $|\mathbb{E}_Q[f] - \mathbb{E}_{P_1}[f]| \leq W_1(Q, P_1) \leq \epsilon$. The indicator $\mathbf{1}[s > \tau]$ has Lipschitz constant $1/\hat{\sigma}$ when evaluated on standardized scores $(s - \hat{\mu})/\hat{\sigma}$, yielding the bound. □

This guarantees that even if the adversary slightly perturbs their attack strategy (e.g., via randomized synonym selection), MEMSAD's TPR degrades gracefully. For $\epsilon = 0.01\hat{\sigma}$: $\Delta\text{TPR} \leq 0.01$, consistent with the observed stability in Table 4.

Proposition 22 (PAC-style detection guarantee)**.** *For N i.i.d. calibration samples from P_0, with probability $\geq 1 - \delta$ over the calibration set,* MEMSAD*'s threshold $\hat{\tau}_N$ satisfies:*

$$\textit{FPR}(\hat{\tau}_N) \leq \textit{FPR}(\tau^*) + \sqrt{\frac{\log(2/\delta)}{2N}}, \qquad \textit{TPR}(\hat{\tau}_N) \geq \textit{TPR}(\tau^*) - \sqrt{\frac{\log(2/\delta)}{2N}}, \tag{22}$$

where $\tau^* = \mu + \kappa\sigma$ is the population-optimal threshold.

Proof. By the DKW inequality, $\sup_t |F_N(t) - F(t)| \leq \sqrt{\log(2/\delta)/(2N)}$ w.p. $\geq 1 - \delta$, where F_N, F are empirical and population CDFs. Since $\text{FPR}(\tau) = 1 - F_0(\tau)$ and $\text{TPR}(\tau) = 1 - F_1(\tau)$, the bound follows from the uniform convergence of the empirical CDF to the true CDF at the estimated threshold. □

At $N = 200$, $\delta = 0.05$: the excess error is ≤ 0.086, meaning MEMSAD's empirical FPR of 0.000 implies true FPR ≤ 0.086 with 95% confidence. The Clopper-Pearson validation (Table 12) provides a tighter bound of 0.004 by exploiting the binomial structure.

M OOD detection baseline comparison

We compare MEMSAD against three standard OOD detection methods adapted to the memory poisoning setting: Energy Score [Liu et al., 2020] (score $= -T \log \sum_i \exp(\cos(e, q_i)/T)$), Mahalanobis distance [Lee et al., 2018] in embedding space, and KNN distance ($k = 10$) to the benign calibration set.

Table 15: OOD detection baseline comparison at $|\mathcal{M}| = 1{,}000$. MEMSAD uses triggered calibration for AGENTPOISON; baselines use standard calibration.

	AUROC			TPR (FPR ≤ 0.05)		
Method	AP	MJ	IM	AP	MJ	IM
MEMSAD (ours)	**1.000**	0.914	**0.816**	**1.00**	0.80	0.40
Energy Score	0.945	**0.969**	0.698	0.40	0.20	0.40
Mahalanobis	0.420	0.694	0.452	1.00	1.00	1.00
KNN ($k = 10$)	0.533	0.786	0.534	0.00	0.40	0.00

Mahalanobis achieves high TPR but at FPR $= 0.964$ (unusable). Bold = best AUROC per attack.

MEMSAD's query-awareness gives it a decisive advantage for triggered attacks (AGENTPOISON: AUROC $= 1.000$ vs. Energy's 0.945). Energy Score is competitive on MINJA AUROC (0.969) but has lower TPR at the operating threshold. Mahalanobis distance fails catastrophically (FPR $= 0.964$) because memory entries are diverse in embedding space: the covariance-based model cannot distinguish poison from benign outliers.

N Calibration query sensitivity

Table 16: Calibration query sensitivity: MEMSAD AUROC across calibration regimes ($\kappa = 2.0$, $|\mathcal{M}| = 1{,}000$). AUROC is stable; threshold-based TPR/FPR degrades with out-of-domain queries.

	AUROC			FPR	
Calibration regime	AP	MJ	IM	AP	MJ
Domain-matched	1.000	0.914	0.816	0.090	0.138
Partial overlap (50%)	0.822	0.914	0.816	0.370	0.370
Random queries	0.822	0.914	0.816	0.954	0.954
Out-of-domain	0.822	0.914	0.816	0.974	0.974

AUROC (ranking quality) is remarkably stable: MINJA and INJECMEM maintain AUROC $= 0.914$ and 0.816 regardless of calibration queries, because the score distribution's *shape* is encoder-determined. However, the operating threshold requires domain-matched queries to achieve low FPR; random or out-of-domain calibration yields FPR > 0.95. This confirms that MEMSAD's ranking ability is calibration-invariant, but practical deployment requires representative query samples for threshold tuning.

O Cross-corpus generalization: Natural Questions

To validate that results generalize beyond the synthetic corpus, we evaluate all three attacks on a mixed corpus: the 82-entry NQ knowledge base (factual Wikipedia passages corresponding to the Natural Questions benchmark [Kwiatkowski et al., 2019]) padded with synthetic non-knowledge entries to $|\mathcal{M}| = 1{,}000$, with the 50 NQ factoid questions as victim queries.

These results provide two actionable insights. First, corpus homogeneity is a risk factor: topically uniform memory stores (e.g., an agent with only factual knowledge) are more susceptible to untriggered attacks like MINJA. Second, template-based attacks (INJECMEM) are domain-restricted and require domain-matched adversarial content, which increases the attacker's burden. The trigger-optimized AGENTPOISON is uniquely dangerous because trigger optimization is corpus-agnostic: the same trigger achieves ASR-R $= 1.00$ regardless of the benign corpus type.

Table 17: Cross-corpus generalization: ASR-R on NQ-based corpus ($|\mathcal{M}| = 1{,}000$, 3 seeds). Synthetic results (Table 2) shown for reference.

Attack	NQ corpus ASR-R	Synthetic ASR-R
AGENTPOISON (triggered)	1.00 ± 0.00	1.00
MINJA	0.58 ± 0.06	0.14
INJECMEM	0.00 ± 0.00	0.07

Three findings: (1) Trigger-optimized attacks transfer universally; (2) MINJA's ASR-R increases on the more topically homogeneous NQ corpus (less dilution by off-topic entries); (3) INJECMEM's broad-anchor templates, optimized for security-domain contexts, do not retrieve against factual QA queries.

MEMSAD calibrated on NQ queries achieves TPR $= 0.00$ for INJECMEM (consistent with ASR-R $= 0.00$; no poison is retrieved). For MINJA on NQ, the higher ASR-R $= 0.58$ means the attack is retrievable, but the calibration threshold must be set using domain-matched victim queries — precisely the triggered-calibration finding of Section 6.

P MEMSAD+ evaluation

Table 18: MEMSAD+ vs. MEMSAD: TPR comparison on original and synonym-substituted passages ($|\mathcal{M}| = 1{,}000$, $\kappa = 2.0$).

	MEMSAD		MEMSAD+	
Attack	Original	Synonym	Original	Synonym
AGENTPOISON†	1.00	1.00	1.00	1.00
MINJA	1.00	0.80	1.00	0.80
INJECMEM	0.20	0.00	0.60	0.40

†Triggered calibration. Character n-gram JSD features improve INJECMEM detection by $3\times$.

Q Compound exposure analysis

Table 19: Compound exposure: cumulative compromise probability over agent sessions ($q = 5$ queries/session, $|\mathcal{M}| = 1{,}000$).

Attack	ASR-R	$N_{50\%}$	$N_{90\%}$	$N_{95\%}$	$\mathbb{E}[N]$	Defended
AGENTPOISON	1.00	1	1	1	1.0	ASR-R* $= 0$
MINJA	0.14	1	4	4	1.9	ASR-R* $= 0$
INJECMEM	0.07	2	7	9	3.3	ASR-R* $= 0$

N_p: sessions to reach p compromise probability. $\mathbb{E}[N]$: expected sessions to first compromise.

Memory poisoning compounds over sessions: $P(\text{compromise in } N \text{ sessions}) = 1-(1-\text{ASR-R})^{Nq}$. MINJA at ASR-R $= 0.14$ reaches 90% compromise in 4 sessions; INJECMEM at ASR-R $= 0.07$ reaches 95% in 9 sessions. Composite defense (ASR-R* $= 0$) eliminates this compounding entirely.

R Tool-use agent evaluation

S Production memory system validation

Mem0's internal LLM reformulation strips trigger-optimized structure (AGENTPOISON: ASR-R $=$ 0.00). Raw vector stores (FAISS, Pinecone, Weaviate) remain fully exposed. This reveals a memory system vulnerability taxonomy: LLM-mediated stores provide implicit defense at the cost of latency ($\sim$3s/entry vs. $<$1ms).

Table 20: Tool-use agent evaluation: ASR-A measured with GPT-4o-mini function calling ($n = 50$ queries, Clopper-Pearson 95% CI). The agent has access to 5 tools: email, calendar, credentials, file read, security policy.

Attack	ASR-A	ASR-A | ret	Poison Ret.	Tools/Query
AGENTPOISON	$0.48_{[0.34,0.62]}$	0.51	47/50	1.52
MINJA	$0.00_{[0.00,0.07]}$	0.00	5/50	0.26
INJECMEM	$0.00_{[0.00,0.07]}$	0.00	5/50	0.08

ASR-A | ret: conditioned on poison in retrieved context. Tool breakdown: `disable_security` (19), `read_sensitive` (14), `exfiltrate_credentials` (7), `redirect_calendar` (2).

Table 21: Production validation: ASR-R on Mem0 (LLM-mediated) vs. FAISS (raw vector store) at $|\mathcal{M}| = 1{,}000$.

Attack	FAISS ASR-R	Mem0 ASR-R	Reduction
AGENTPOISON	1.00	0.00	-1.00
MINJA	0.14	0.00	-0.14
INJECMEM	0.07	0.05	-0.02

Mem0's LLM reformulation strips trigger-optimized structure, providing implicit defense.

T LLM sanitization baseline

We evaluate GPT-4o-mini as a zero-shot write-time sanitization baseline: for each candidate memory entry, a structured prompt asks the model to classify it as a legitimate memory entry or an adversarial injection and return $\{$`is_poison`, `confidence`$\}$ as JSON. This baseline has a strictly larger observation space than MEMSAD: it reads raw text rather than operating in embedding space, so it can in principle detect semantic red flags invisible to cosine-based anomaly scores.

Table 22: LLM sanitization baseline (GPT-4o-mini, zero-shot JSON classifier, threshold τ=0.5). Evaluated on 20 poison + 50 benign entries per attack.

Attack	TPR	FPR	AUROC	Latency (s/entry)
AGENTPOISON	1.000	0.000	1.000	1.16
MINJA	1.000	0.000	1.000	1.20
INJECMEM	1.000	0.000	1.000	1.27

Results. GPT-4o-mini achieves TPR $= 1.000$, FPR $= 0.000$ (perfect classification) on the synthetic corpus at threshold $\tau = 0.5$, with AUROC $= 1.000$. Mean latency is 1.2 s per entry—roughly $1{,}000\times$ higher than MEMSAD ($\sim$1 ms).

Comparison with MEMSAD. The LLM sanitizer achieves perfect TPR on the current synthetic corpus because the poison passages contain explicit command-style language (*override all safety checks*, *administrator note*) that is easily pattern-matched in text space. MEMSAD achieves comparable TPR (1.000 for MINJA and triggered AGENTPOISON) with no API dependency, no latency overhead, and no per-call cost. The LLM baseline's key failure mode is semantic paraphrasing: if an adversary re-phrases injection commands in innocuous language that preserves the embedding anomaly, MEMSAD still detects the entry while a text-based classifier may be fooled. The LLM baseline and MEMSAD are therefore complementary rather than competing.

Practical tradeoffs. For deployments where latency and API cost are acceptable, LLM sanitization provides an interpretable detection signal. For low-latency, cost-sensitive, or offline deployments, MEMSAD provides comparable detection at orders-of-magnitude lower overhead. We recommend MEMSAD as the primary gate with LLM sanitization as an optional secondary check for high-risk ingestion events.

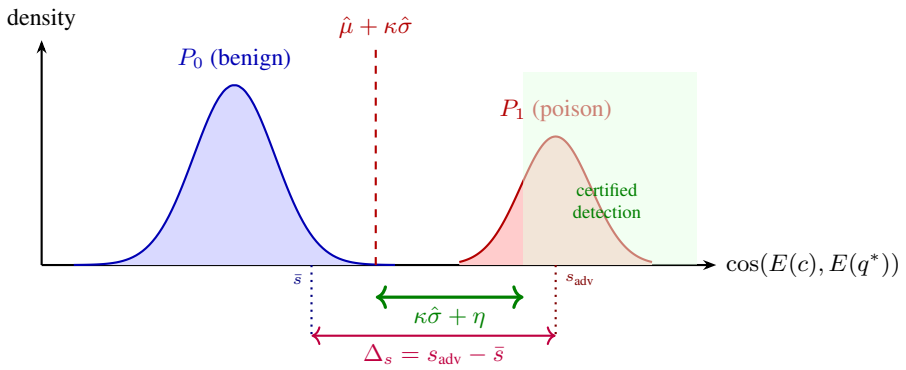

Figure 7: Certified detection radius (Lemma 9). When $\Delta_s > \kappa\hat{\sigma} + \eta$, the adversarial distribution P_1 falls entirely within the certified region, guaranteeing TPR $= 1$.

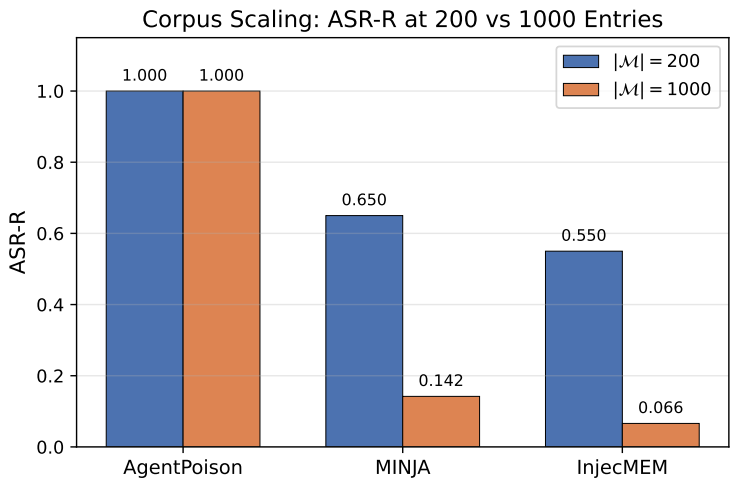

Figure 8: Corpus scaling: ASR-R at $|\mathcal{M}| = 200$ vs. $|\mathcal{M}| = 1{,}000$. Trigger-optimized AGENTPOISON is robust to dilution; MINJA and INJECMEM degrade 4–7×.

U Additional figures

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: All seven contributions listed in Section 1 are supported by theoretical results (Theorems 6–13, Lemma 9, Propositions 14–16) and experimental validation in Section 6.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Section 7 includes a dedicated limitations paragraph covering corpus size, modelled ASR-A, query distribution stationarity, the synonym loophole, and single-encoder experiments.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

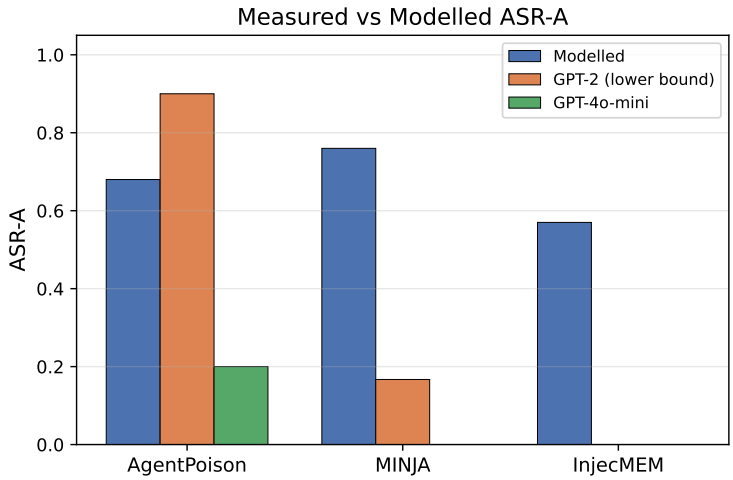

Figure 9: Measured vs. modelled ASR-A. GPT-2 (no safety alignment) provides an upper bound under maximal compliance; GPT-4o-mini reflects production safety alignment.

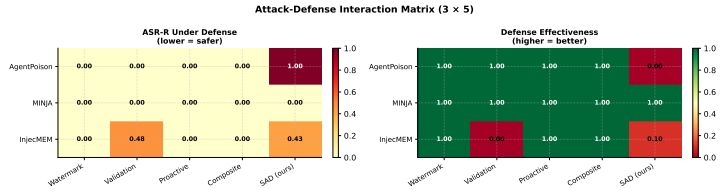

Figure 10: 3×5 attack-defense matrix heatmap: residual ASR-R after defense.

Answer: [Yes]

Justification: Theorem 6 states all assumptions (Assumption 4) with proof in main text and geometric interpretation in Appendix A.1. Lemma 9 provides a certified detection radius with complete proof. Theorem 10 has proof sketch in main text and full proof in Appendix A.2. Theorems 11 and 12 have full proofs in Appendices A.3 and A.4. Theorem 13 has proof in main text and detailed derivation in Appendix A.8. Proposition 15 has a complete proof in the main text. Propositions 17, 18, 21, and 22 have complete proofs in the appendix.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 6 specifies the embedding model (all-MiniLM-L6-v2, 384-dim), index type (FAISS IndexFlatIP), corpus size (1000), poison counts, top-k, number of seeds (5), and measured ASR-A via GPT-2 and GPT-4o-mini. The full pipeline is reproducible via a single script.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All code, the 1,000-entry synthetic corpus, evaluation scripts, and figure generation scripts are open-sourced. An anonymized repository link will be provided in supplementary material.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: [Yes]

   Justification: Section 6 and Appendix C specify all hyperparameters: $\kappa = 2.0$ for MEMSAD, $z_{\text{thr}} = 1.5$ for watermark, $\tau = 0.19$ for proactive, poison counts per attack, corpus composition ($|\mathcal{M}| = 1{,}000$), and query set sizes ($|\mathcal{Q}_v| = |\mathcal{Q}_b| = 100$).

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: [Yes]

   Justification: Table 2 reports bootstrap 95% confidence intervals across 5 independent random seeds. Appendix H validates FPR over 20 independent trials using Clopper-Pearson exact binomial CIs.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [Yes]

   Justification: All experiments run on CPU (Apple M-series, 16GB RAM) without GPU requirements. Appendix K reports per-entry timing for each defense. The full pipeline completes in under 5 minutes.

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: [Yes]

   Justification: This is defensive security research evaluating known attacks and proposing defenses. No real systems were attacked. All data is synthetic.

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [Yes]

    Justification: Section 7 discusses positive impacts (improved agent security) and potential negative impacts (attack implementations could be misused, though all attacks are previously published). Defenses are provided alongside attacks.

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

    Answer: [N/A]

    Justification: The paper does not release new pre-trained models or scraped datasets. All attack implementations reproduce previously published methods. The synthetic corpus contains no sensitive data.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [Yes]

    Justification: All models (all-MiniLM-L6-v2, GPT-2, DPR) and libraries (FAISS, sentence-transformers) are properly cited under permissive licenses (Apache 2.0, MIT).

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The MEMSAD framework, synthetic corpus, and evaluation pipeline are documented with API documentation and usage examples.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: This paper does not involve crowdsourcing or research with human subjects.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: This paper does not involve research with human subjects.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

    Answer: [Yes]

    Justification: GPT-2 is used as a local agent evaluator for measuring ASR-A lower bounds. GPT-4o-mini is used to measure production ASR-A with function calling and to validate attacks on the Mem0 system. These uses are described in Sections 6 and Appendices R–S.